

We would like to thank the four reviewers for their feedback. Upon acceptance, we will include in the final version (a) a clearer presentation of the numerical results and (b) missing references. We first discuss a common concern shared by reviewer 1, reviewer 2, reviewer 4.

Novelty of The Contribution: We want to stress on the generality of our incremental framework, which tackles a constrained, non-convex and non-smooth optimization problem. The main contribution of this paper is to propose a unifying framework for the analysis of a large class of optimization algorithms which indeed includes well-known but not so well-studied algorithms. The major goal here is to relax the class of surrogate functions used in MISO [Mairal, 2015] and replace them by the respective Monte-Carlo approximations. We provide a general algorithm and global convergence analysis under mild assumptions on the model and show that two examples, MLE for general latent data models and Variational Inference, are its special instances. Working at the crossroads of *Optimization* and *Sampling* constitutes what we believe to be the novelty and the technicality of our theoretical results.

Reviewer 1: We thank the reviewer for valuable comments and references. We would like to make the following clarification regarding the difference with MISO:

Originality: The main contribution of the paper is to extend the MISO algorithm when the surrogate functions are not tractable. We motivate the need for dealing with intractable surrogate functions when nonconvex latent data models are being trained. In this case, the surrogate functions can be written as an expectation due to the latent structure of the problem and the nonconvexity yields a generally intractable expectation to compute. The only option is to build a stochastic surrogate function based on a Monte Carlo approximation.

Reviewer 2: We thank the reviewer for the useful comments. Our point-to-point response is as follows:

Numerical Plots: Due to space constraints, we only presented several dimensions for the logistic parameters and the mean of the latent variable. As the reviewer mentioned, we also learn the variance of those latent variables and the convergence plots of those variances will be added to the rebuttal version. For all experiments, we made sure that the estimated parameters for each method converge to the same value. This was our main criteria to claim that a method is faster than the other. Then, the problem being (highly) nonconvex, the estimations can indeed get trapped in various local minima. Regardless of generalization properties of the output vector of estimated parameters, our focus through those numerical examples was to highlight faster convergence, in iteration, of our method.

Wallclock Time: Wallclock time per iteration is comparable for each method. Indeed the methods always only involve first order computation. Running times table will be provided in the revised paper.

Parameter Tuning: The baseline methods were tuned and presented to the best of their performances both with regards to their stepsize (grid search) and minibatch size. We believe your remark refers to the first numerical example (logistic regression with missing values): Regarding the stepsize, as MCEM does not have one, we indeed tuned the stepsize of SAEM. Rather than c/k , common practice is to tune a parameter α such that $\gamma_k = 1/\gamma^\alpha$. We report results for SAEM with the best α ($\alpha = 0.6$). Regarding batch size, for SAEM and MCEM both are full batch methods and the idea here is to compare different values of minibatch size for the MISSO method to see its influence on the performances.

Reviewer 3: We thank the reviewer for valuable comments and references. We clarify the following point:

Verification of the Assumptions: For the Logistic regression with missing values, we consider that the covariance matrix is PSD and make sure in practice that its smallest eigenvalue is away from zero. For the Variational inference example, we recall the reviewer that the updates Section C.2 stem from the minimization of the quadratic functions Eq.(11) and that indeed a projection step needs to be added in order to ensure boundedness of the iterates. For illustrative purposes we did not implement the algorithm that sticks closely to our formulation but after careful consideration, the surrogate problem being a convex one on a convex closed set, the implementation will become more complex without obstructing the theory. We will provide the two variants of the algorithm in the rebuttal.

Reviewer 4: We thank the reviewer for valuable comments and references. Our point-to-point response is as follows:

Comparison to [Murray+, 2012] and [Tran+, 2017]: [Murray+, 2012] is out of scope of our paper since their focus is on purely Bayesian models where the normalizing constant depends on the latent variable. In such case, since you cannot run standard MH algorithm (normalizing constants do not cancel out in the MH ratio), the authors develop a new MCMC method. [Tran+, 2017] is relevant to our paper and will be included in the rebuttal. Though, their framework is only a *full-batch* instance of our general MISSO scheme which includes incremental Variational Inference (ELBO maximization, see Example 2) but also missing values problem which is a totally different setting.

Comparison to [Kang+, 2015]: [Kang+, 2015] solely focuses on MM scheme when the surrogate functions are deterministic, *i.e.* can be computed exactly and using full batch update (versus our incremental and scalable update). Also, their analysis requires *strong convexity* of the gap between the convex surrogate and the nonconvex objective function while our analysis only requires a *smoothness* assumption, see H2.