Convergent Adaptive Gradient Methods in Decentralized Optimization

Anonymous Author(s)
Affiliation
Address
email

Abstract

Adaptive gradient methods including Adam, AdaGrad, and their variants have been very successful for training deep learning models, such as neural networks, in the past few years. Meanwhile, given the need for distributed training procedure, the realm of distributed optimization algorithms is at the center of attention. With the growth of computing power and the need for using machine learning models on mobile devices, the communication cost of distributed training algorithms needs careful consideration. In response to this, more and more attention is shifted from the traditional parameter server training paradigm to the decentralized one, which usually requires lower communication costs. In this paper, we try to rigorously incorporate adaptive gradient methods into decentralized training procedures, coming up with convergent decentralized adaptive gradient methods. Specifically, we propose a general algorithmic framework that can convert existing adaptive gradient methods to their decentralized counterparts. In addition, we rigorously analyze the convergence behavior of the proposed algorithmic framework and show that if a given adaptive gradient method converges, under some specific conditions, then its decentralized counterpart is also convergent.

1 Introduction

Distributed training of machine learning models is drawing growing attention in the past few years due to its practical benefits and necessities. Given the evolution of computing capabilities of CPUs and GPUs, computation time in distributed setting is gradually dominated by the communication time in many circumstances [Chilimbi et al., 2014, McMahan et al., 2016]. As a result, a large amount of recent works has been focussing on reducing communication cost for distributed learning [Alistarh et al., 2017, Lin et al., 2017, Wangni et al., 2018, Stich et al., 2018, Wang et al., 2018, Tang et al., 2019]. In the traditional parameter (central) server setting, where a parameter server is employed to manage communication in the whole network, many effective communication reductions have been proposed based on gradient compression [Aji and Heafield, 2017] and quantization [Chen et al., 2010, Ge et al., 2013, Jegou et al., 2010]. Despite these communication reduction techniques, its cost still, usually, scales linearly with the number of workers. Due to this limitation and with the sheer size of decentralized devices, the *decentralized training paradigm* [Duchi et al., 2011b], where the parameter server is removed and each node only communicates with its neighbors, is drawing attention. It has been shown in Lian et al. [2017] that decentralized training algorithms can outperform parameter server-based algorithms when the training bottleneck is the communication cost. The decentralized training paradigm is also naturally preferred when a parameter server is not available.

In light of recent advances in nonconvex optimization, an effective way to accelerate training is by using adaptive gradient methods like AdaGrad [Duchi et al., 2011a], Adam [Kingma and Ba, 2014] or AMSGrad [Reddi et al., 2019]. Their practical benefits are proven by their popularity in training neural networks, featured by faster convergence and ease of parameter tuning compared with SGD. Despite a large amount of literature in distributed optimization, there have been few

Submitted to 34th Conference on Neural Information Processing Systems (NeurIPS 2020). Do not distribute.

works considering bringing adaptive gradient methods into distributed training, largely due to the lack of understanding of adaptive gradient methods convergence behaviors. Notably, Reddi et al. [2020] develop the first decentralized ADAM method for distributed optimization problems with a direct application to federated learning. An inner loop is employed to computed mini-batch gradients on each worker nodes and a global adaptive step is done to update the global parameter at each central-server iteration. Yet, in the settings of our paper, nodes can only communicate with their neighbors while a server/worker communication is needed in [Reddi et al., 2020]. Designing adaptive methods in such settings is highly non-trivial due to the already complicated update rules and the interaction between the effect of using adaptive learning rates and the decentralized communication protocols.

This paper is an attempts at bridging the gap between both realms in nonconvex optimization. Our contributions are summarized as follows:

- In this paper, we investigate the possibility of using any adaptive gradient methods in the decentralized training paradigm. We develop a general technique that can convert an adaptive gradient method from a centralized method to a decentralized method.
- By using our proposed technique, we present a new decentralized optimization algorithm, called decentralized AMSGrad, as the decentralized counterpart of AMSGrad.
- We provide a theoretical verification interface for analyzing the behavior of decentralized adaptive gradient methods obtained as a result of our technique. Built upon our proposed analysis framework for that type of decentralized algorithms, we can characterize the convergence rate of decentralized AMSGrad, which is the first convergent decentralized adaptive gradient method.

A *novel technique* in our framework is a mechanism to enforce a consensus on adaptive learning rates at different nodes. We show the importance of consensus on adaptive learning rates by proving a divergent problem instance for a recently proposed decentralized adaptive gradient method DADAM [Nazari et al., 2019], a decentralized version of ADAM, which lacks consensus mechanisms on adaptive learning rates.

After presenting related work and important concepts of decentralized adaptive methods in Section 2, we develop our general framework for converting any adaptive gradient algorithm in its decentralized counterpart along with their rigorous finite-time convergence analysis in Section 3. Section 4 and 5 conclude our work with illustrative examples of our framework.

Notations: $x_{t,i}$ denotes variable x at node i and iteration t. $\|\cdot\|_{abs}$ denotes the entry-wise L_1 norm of a matrix, i.e. $\|A\|_{abs} = \sum_{i,j} A_{i,j}$. We introduce important notations used throughout the paper: for any $t >$, $G_t := [g_{t,N}]$ where $[g_{t,N}]$ denotes the vector $[g_{t,1}, g_{t,2}, \cdots, g_{t,N}]$, $M_t := [m_{t,N}]$, $X_t := [x_{t,N}]$, $\overline{\nabla f}(X_t) := \frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i})$, $U_t := [u_{t,N}]$, $\tilde{U}_t := [\tilde{u}_{t,N}]$, $V_t := [v_{t,N}]$, $\hat{V}_t := [\hat{v}_{t,N}]$, $\overline{X}_t := \frac{1}{N}\sum_{i=1}^N x_{t,i}$, $\overline{U}_t := \frac{1}{N}\sum_{i=1}^N u_{t,i}$ and $\overline{\tilde{U}_t} := \frac{1}{N}\sum_{i=1}^N \tilde{u}_{t,i}$.

2 Decentralized Adaptive Training and Divergence of DADAM

2.1 Related Work

Decentralized optimization: Traditional decentralized optimization methods include well-know algorithms such as ADMM [Boyd et al., 2011], dual averaging [Duchi et al., 2011b], distributed subgradient descent [Nedic and Ozdaglar, 2009]. More recent algorithms include Extra [Shi et al., 2015], Next [Di Lorenzo and Scutari, 2016] and Prox-PDA [Hong et al., 2017]. While these algorithms were commonly used in applications other than deep learning, recent algorithmic advances in the machine learning community have shown that decentralized optimization can be useful for training deep models such as neural networks too. Lian et al. [2017] show that a stochastic version of decentralized subgradient descent can outperform parameter server-based algorithms when the communication cost is high. Tang et al. [2018] propose the D^2 algorithm improving the convergence rate over stochastic subgradient descent. Assran et al. [2018] propose the Stochastic Gradient Push that is more robust to network failures for training neural networks. The study of decentralized training in the machine learning community is only at its initial stage. No one has seriously considered designing adaptive gradient methods in the setting of decentralized training until a recent work [Nazari

et al., 2019] where a decentralized version of AMSGrad [Reddi et al., 2019] is proposed is proven to satisfy some non-standard regret.

Adaptive gradient methods: Adaptive gradient methods have been popular in recent years due to their superior performance in training neural networks. Most used adaptive methods include AdaGrad [Duchi et al., 2011a] or Adam [Kingma and Ba, 2014] and their variants. Key features of such methods lie in the use of momentum and adaptive learning rates (which means the learning rate is changing during optimization and are anisotropic, i.e. depend on the dimension). The method of reference, Adam, has been analyzed in [Reddi et al., 2019] where the authors point out an error in previous convergence analyses. Since then, a range of works have been focussing on analyzing the convergence behavior of the various existing adaptive gradient methods. Ward et al. [2018], Li and Orabona [2018] derive convergence guarantees for a variant of AdaGrad without coordinate-wise learning rates. Chen et al. [2018] analyze the convergence behavior of a broad class of algorithms including AMSGrad [Reddi et al., 2019] and AdaGrad. Zou and Shen [2018] provide a unified convergence analysis for AdaGrad with momentum. A few recent adaptive gradient methods can be found in [Agarwal et al., 2018, Luo et al., 2019, Zaheer et al., 2018].

2.2 Decentralized Optimization

In distributed optimization (with N nodes), we aim at solving the following problem

$$\min_{x\in\mathbb{R}^d} \frac{1}{N}\sum_{i=1}^{N} f_i(x)\,, \tag{1}$$

where x is the vector of parameters and f_i is only accessible by the ith node. Through the prism of neural network training procedure, f_i can be viewed as the average loss of the data samples located at node i. Throughout the paper, we make the following assumptions for analyzing the convergence behavior of the different algorithms.

A1. *For all $i \in [N]$, f_i is differentiable and the gradients is L-Lipschitz, i.e. $\|\nabla f_i(x) - \nabla f_i(y)\| \leq L\|x-y\|, \forall x, y$.*

A2. *We assume at iteration t, node i can access a stochastic gradient $g_{t,i}$. In addition, the stochastic gradients have bounded L_∞ norm and the gradients of f_i are also bounded, i.e. $\|g_{t,i}\| \leq G_\infty$, $\|\nabla f_i(x)\|_\infty \leq G_\infty$.*

A3. *The gradient estimators are unbiased and each coordinate have bounded variance, i.e. $\mathbb{E}[g_{t,i}] = \nabla f_i(x_{t,i})$ and $\mathbb{E}[([g_{t,i} - f_i(x_{t,i})]_j)^2] \leq \sigma^2, \forall t, i, j$.*

The assumptions A1 and A3 are standard in distributed optimization. A2 is slightly stronger than the traditional assumption that the estimator has bounded variance, but is commonly used for the analysis of adaptive gradient methods [Chen et al., 2018, Ward et al., 2018]. One thing that should be noted is that the bounded gradient estimator assumption in A2 implies the bounded variance assumption in A3. We denote the variance bound and the estimator bound differently to avoid confusion when we use them for different purposes. In decentralized optimization, the nodes are connected as a graph and each node only communicates to its neighbors. In such cases, one usually constructs a $N \times N$ matrix W for information sharing when designing new algorithms. We denote λ_i to be its ith largest eigenvalue and define $\lambda \triangleq \max(|\lambda_2|, |\lambda_N|)$. As can be expected, W cannot be arbitrary, the key properties required for W are listed in A4.

A4. *The matrix W satisfies:* (I) $\sum_{j=1}^{N} W_{i,j} = 1$, $\sum_{i=1}^{N} W_{i,j} = 1$, $W_{i,j} \geq 0$, (II) $\lambda_1 = 1$, $|\lambda_2| < 1$, $|\lambda_N| < 1$ *and* (III) $W_{i,j} = 0$ *if node i and node j are not neighbors.*

2.3 Divergence of DADAM

Recently, Nazari et al. [2019] initiated a trial to bring adaptive gradient methods into decentralized optimization, the resulting algorithm is DADAM, which is shown in Algorithm 1. DADAM is essentially a decentralized version of AMSGrad and the key modification is the use of a consensus step on optimization variable x to transmit information across the network, encouraging convergence. The matrix W is a doubly stochastic matrix (which satisfies A4) for achieving average consensus of x. Introducing such mixing matrix is a standard approach for decentralizing an algorithm, such as distributed gradient descent [Nedic and Ozdaglar, 2009, Yuan et al., 2016]. It is proven in Nazari et al. [2019] that DADAM admits a non-standard regret bound in the online setting, however, whether the algorithm can converge to stationary points in standard offline settings such training neural networks is still unknown.

In the following, we show the DADAM may fail to converge in the offline nonconvex optimization settings.

Theorem 1. *There exist a problem satisfying A1 − A4 where DADAM fail to converge.*

Proof. Consider a 1 dimensional optimization problem distributed on two nodes $\min_x \frac{1}{2}\sum_{i=1}^{2} f_i(x)$ where $f_i(x) = \frac{1}{2}(x - a_i)^2$ and $a_1 = 0$, $a_2 = 1$. The network contains only two nodes and the matrix W satisfies $W_{ij} = \frac{1}{2}$ for all i, j. For simplicity, we consider running DADAM with $\beta_1 = \beta_2 = \beta_3 = 0$ and $\epsilon = 0.6$. Suppose we initialize DADAM at $x_{1,i} = 0$ for all $i \in [N]$ and use the following learning rate $\alpha = 0.001$. We have at $x_{1,i} = 0$, $\nabla f_1(x_{1,1}) = 0$, $\nabla f_2(x_{1,2}) = 1$, leading to $\hat{v}_{1,1} = 0.6$ and $\hat{v}_{1,2} = 1$. Thus, from step 1, we will have $\hat{v}_{1,2} \geq 1$. In addition, it is can be easily proved that, with the stepsize selection, we always have $\hat{v}_{1,1} < 1$, in fact, it will not reach 0.6. Thus, in the next iterations, the gradient of losses on node 1 and 2 will be scaled differently. This scaling is equivalent to running gradient descent on a objective where the losses of the two nodes are scaled by different factors. In such case, the algorithm will converge to a stationary point of a weighted average of the loss on node 1. Recall that the problem we tackle to illustrate Theorem 1 is a quadratic problem with only one minimizer. Then, since the weight of the losses on the two nodes are different and that the unbalanced weights on the two functions yields a different minimizer, the algorithm will not converge to the unique stationary point of the original loss (which is $x = 0.5$). □

Algorithm 1 DADAM (with N nodes)

Input: α, current point X_t, $u_{\frac{1}{2},i} = \hat{v}_{0,i} = \epsilon\mathbf{1}$, $m_0 = 0$ and mixing matrix W
for $t = 1, 2, \cdots, T$ **do**
  for all $i \in [N]$ **do in parallel**
    $g_{t,i} \leftarrow \nabla f_i(x_{t,i}) + \xi_{t,i}$
    $m_{t,i} = \beta_1 m_{t-1,i} + (1 - \beta_1) g_{t,i}$
    $v_{t,i} = \beta_2 v_{t-1,i} + (1 - \beta_2) g_{t,i}^2$
    $\hat{v}_{t,i} = \beta_3 \hat{v}_{t,i} + (1 - \beta_3) \max(\hat{v}_{t-1,i}, v_{t,i})$
    $x_{t+\frac{1}{2},i} = \sum_{j=1}^{N} W_{ij} x_{t,j}$
    $x_{t+1,i} = x_{t+\frac{1}{2},i} - \alpha \frac{m_{t,i}}{\sqrt{\hat{v}_{t,i}}}$
end for

Theorem 1 claims that even though DADAM is proven to satisfy some regret bounds , see [Nazari et al., 2019], it can fail to converge to stationary points in the nonconvex offline setting, which is a common setting for training neural networks. We conjecture that this inconsistency is due to the definition of the regret in [Nazari et al., 2019]. In the next section, we will design decentralized adaptive gradient methods that are guaranteed to converge to stationary points and provide a characterization of that convergence in finite-time and independently of the initialization of our methods.

3 Decentralized Adaptive Gradient Methods and their Convergence

In this section, we discuss the difficulties of designing adaptive gradient methods in decentralized optimization and introduce an algorithmic framework that convert existing convergent adaptive gradient methods to their decentralized counterparts. We also develop the first convergent decentralized adaptive gradient method, converted from AMSGrad, as an instance of this proposed framework.

3.1 Importance and Difficulties of Consensus on Adaptive Learning Rates

The divergent example in the previous section implies that we should synchronize the adaptive learning rates on different nodes. This can be easily achieved in the parameter server setting where all the nodes are sending their gradients to a central server at each iteration. The parameter server can then exploit the received gradients to maintain a sequence of synchronized adaptive learning rates when updating the parameters, see [Reddi et al., 2020]. However, in our setting of decentralized training, every node can only communication with its neighbors and such central parameter server does not exist. Under that setting, the information for updating the adaptive learning rates can be only shared locally instead of broadcasted over the whole network. This makes it impossible to obtain, in a single iteration, a synchronized adaptive learning rate update using all the information in the network.

Systemic Approach: On a systemic level, one way to alleviate this bottleneck is to design communication protocols to give each node access to the same aggregated gradients over the whole network at least periodically if not at every iteration. Therefore, the nodes can update their individual adaptive learning rates based on the same information. However, such solution introduce extra communication cost since it involves broadcasting over the network.

Algorithmic Approach: Our contributions being on an algorithmic level, another way to solve the aforementioned problem is by letting the sequences of adaptive learning rates, present on different

Algorithm 2 Decentralized Adaptive Gradient Method (with N nodes)

Input: learning rate α, initial point $x_{1,i} = x_{init}, u_{\frac{1}{2},i} = \hat{v}_{0,i}, m_{0,i} = 0$, mixing matrix W
for $t = 1, 2, \cdots, T$ **do**
  for all $i \in [N]$ **do in parallel**
    $g_{t,i} \leftarrow \nabla f_i(x_{t,i}) + \xi_{t,i}$
    $m_{t,i} = \beta_1 m_{t-1,i} + (1-\beta_1) g_{t,i}$
    $\hat{v}_{t,i} = r_t(g_{1,i}, \cdots, g_{t-1,i})$
    $x_{t+\frac{1}{2},i} = \sum_{j=1}^{N} W_{ij} x_{t,j}$
    $\tilde{u}_{t,i} = \sum_{j=1}^{N} W_{ij} \tilde{u}_{t-\frac{1}{2},j}$
    $u_{t,i} = \max(\tilde{u}_{t,i}, \epsilon)$
    $x_{t+1,i} = x_{t+\frac{1}{2},i} - \alpha \frac{m_{t,i}}{\sqrt{u_{t,i}}}$
    $\tilde{u}_{t+\frac{1}{2},i} = \tilde{u}_{t,i} - \hat{v}_{t-1,i} + \hat{v}_{t,i}$
end for

nodes, to *consent* gradually, through the iterations. Intuitively, if the adaptive learning rates can consent fast enough, the difference among the adaptive learning rates on different nodes will not affect the convergence of the algorithm. The benefit of such approach is that we do not need to introduce extra communication cost.

3.2 Decentralized Adaptive Gradient Unifying Framework

As mentioned before, we need to choose a method to implement consensus of adaptive learning rates. While each node can have different $\hat{v}_{t,i}$ in DADAM, one can keep track of the min/max/average of these adaptive learning rates and use this quantity to update the adaptive learning rates. Also one can predefine some convergent lower and upper bounds to gradually synchronize the adaptive learning rates on different nodes as developed for AdaBound in [Luo et al., 2019]. In this paper, we opt for the average consensus on $\hat{v}_{t,i}$. Since in adaptive gradient methods such as AdaGrad or Adam, the quantity $\hat{v}_{t,i}$ approximates the second moment of the gradient estimator, the average of the estimations of those second moments from different nodes is an estimation of second moment on the whole network. Also, this design will not introduce any extra hyperparameters that can potentially complicate the tuning process. We now present the main convergence result for our class methods:

Theorem 2. *Assume A1-A4. Set* $\alpha = 1/\sqrt{Td}$. *When* $\alpha \leq \frac{\epsilon^{0.5}}{16L}$, *Algorithm 2 yields the following regret bound*

$$\frac{1}{T}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right] \leq C_1\frac{\sqrt{d}}{\sqrt{T}}\left(\mathbb{E}[f(Z_1)] - \min_z f(z) + \frac{\sigma^2}{N}\right) + \frac{C_2}{T} + \frac{C_3}{T^{1.5}d^{0.5}} + \left(\frac{C_4}{TN^{0.5}} + \frac{C_5}{T^{1.5}d^{0.5}N^{0.5}}\right)\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs}\right], \quad (2)$$

where $\|\cdot\|_{abs}$ *denotes the entry-wise* L_1 *norm of a matrix (i.e* $\|A\|_{abs} = \sum_{i,j}|A_{ij}|$*). The constants* $C_1 = \max(4, 4L/\epsilon)$, $C_2 = 6((\beta_1/(1-\beta_1))^2 + 1/(1-\lambda)^2)LG_\infty^2/\epsilon^{1.5}$, $C_3 = 16L^2(1-\lambda)G_\infty^2/\epsilon^2$, $C_4 = 2/(\epsilon^{1.5}(1-\lambda))(\lambda+\beta_1/(1-\beta_1))G_\infty^2$, $C_5 = 2/(\epsilon^2(1-\lambda))L(\lambda+\beta_1/(1-\beta_1))G_\infty^2 + 4/(\epsilon^2(1-\lambda))LG_\infty^2$ *are independent of d, T and N. In addition,* $\frac{1}{N}\sum_{i=1}^{N}\left\|x_{t,i} - \overline{X}_t\right\|^2 \leq \alpha^2\left(\frac{1}{1-\lambda}\right)^2 dG_\infty^2\frac{1}{\epsilon}$ *which quantifies the consensus error.*

Theorem 2 shows that if $\mathbb{E}[\sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs}] = o(T)$ and $\bar{U}_t$ is upper bounded, then Algorithm 2 is guaranteed to converge to stationary points of the regret function. Intuitively, this means that if the adaptive learning rates on different nodes do not change too fast, the algorithm can converge. This is true as stated in [Chen et al., 2018] where it is shown that if such condition is violated, the algorithm can diverge. Furthermore, Theorem 2 conveys the benefits of using more nodes. As N becomes larger, the term σ^2/N will be small. This is also strengthened by the fact that with the growth of N, the training process tends to be more stable. We now present, in Algorithm 3, a notable special case of our algorithmic framework, namely Decentralized AMSGrad, which is a decentralized variant of AMSGrad.

Compared with DADAM, the above algorithm leverages a dynamic average consensus mechanism to keep track of average of $\{\hat{v}_{t,i}\}_{i=1}^N$, stored as $\tilde{u}_{t,i}$ on ith node, and uses $u_{t,i} = \max(\tilde{u}_{t,i}, \epsilon)$ for updating the adaptive learning rate for ith node. As the number of iteration grows, even though $\hat{v}_{t,i}$ on different nodes can converge to different constants, all the $u_{t,i}$ will be converge to the same number $\lim_{t\to\infty} 1/N \sum_{i=1}^N \hat{v}_{t,i}$ if the limit exists. The use of this average consensus mechanism enables the consensus of adaptive learning rates on different nodes, which consequentially guarantees convergence to stationary points. The consensus of adaptive learning rates is the key difference between decentralized AMSGrad and DADAM and is the reason why decentralized AMSGrad is a convergent algorithm while DADAM is not. One may noticed that decentralized AMSGrad does not deduce to AMSGrad because $u_{t,i}$ in line 10 is calculated based on $v_{t-1,i}$ instead of $v_{t,i}$. This encourages parallel execution of gradient computation and communication. Specifically, line 4-7 in Algorithm 3 and Algorithm 2 can be executed in parallel with line 8-9 to overlap communication and computation time. If $u_{t,i}$ depends on $v_{t,i}$ which in turn depends on $g_{t,i}$, the gradient computation must finish before the consensus step of adaptive learning rate line in 9. This can slow down per-iteration running time of the algorithm. To avoid such delayed adaptive learning, adding $\tilde{u}_{t-\frac{1}{2},i} = \tilde{u}_{t,i} - \hat{v}_{t-1,i} + \hat{v}_{t,i}$ before line 9 and get rid of line 12 in Algorithm 2 is an option. Similar convergence guarantees will hold since one can easily modify our proof of Theorem 2 for such update rule. As stated above, Algorithm 3 converges, with the following rate:

Algorithm 3 Decentralized AMSGrad (with N nodes)

1: **Input:** learning rate α, initial point $x_{1,i} = x_{init}, u_{\frac{1}{2},i} = \hat{v}_{0,i} = \epsilon\mathbf{1}$ (with $\epsilon \geq 0$), $m_{0,i} = 0$, mixing matrix W
2: **for** $t = 1, 2, \cdots, T$ **do**
3: **for all** $i \in [N]$ **do in parallel**
4: $g_{t,i} \leftarrow \nabla f_i(x_{t,i}) + \xi_{t,i}$
5: $m_{t,i} = \beta_1 m_{t-1,i} + (1-\beta_1) g_{t,i}$
6: $v_{t,i} = \beta_2 v_{t-1,i} + (1-\beta_2) g_{t,i}^2$
7: $\hat{v}_{t,i} = \max(\hat{v}_{t-1,i}, v_{t,i})$
8: $x_{t+\frac{1}{2},i} = \sum_{j=1}^N W_{ij} x_{t,j}$
9: $\tilde{u}_{t,i} = \sum_{j=1}^N W_{ij} \tilde{u}_{t-\frac{1}{2},j}$
10: $u_{t,i} = \max(\tilde{u}_{t,i}, \epsilon)$
11: $x_{t+1,i} = x_{t+\frac{1}{2},i} - \alpha \frac{m_{t,i}}{\sqrt{u_{t,i}}}$
12: $\tilde{u}_{t+\frac{1}{2},i} = \tilde{u}_{t,i} - \hat{v}_{t-1,i} + \hat{v}_{t,i}$
13: **end for**

Theorem 3. *Assume A1-A4. Set $\alpha = 1/\sqrt{Td}$. When $\alpha \leq \frac{\epsilon^{0.5}}{16L}$, Algorithm 3 yields the following regret bound*

$$\frac{1}{T}\sum_{t=1}^T \mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right] \leq C_1' \frac{\sqrt{d}}{\sqrt{T}}\left(\mathbb{E}[f(Z_1)] - \min_z f(z) + \frac{\sigma^2}{N}\right) + \frac{C_2'}{T} + \frac{d}{T}\sqrt{N}C_4' + \frac{\sqrt{d}}{T^{1.5}}\sqrt{N}C_5',$$

where $C_1' = C_1$, $C_2' = C_2$, $C_3' = C_3$, $C_4' = C_4 G_\infty^2$ and $C_5' = C_5 G_\infty^2$. C_1, C_2, C_3, C_4, C_5 are constants independent of d, T and N defined in Theorem 2. In addition, the consensus of variables at different nodes is given by $\frac{1}{N}\sum_{i=1}^N \left\|x_{t,i} - \overline{X}_t\right\|^2 \leq \frac{1}{T}\left(\frac{1}{1-\lambda}\right)^2 G_\infty^2 \frac{1}{\epsilon}$.

Theorem 3 shows that Algorithm 3 converges with a rate of $\mathcal{O}\sqrt{d}/\sqrt{T})$ when T is large, which is the best known convergence rate under the given assumptions. Note that in some related works, SGD admits a convergence rate of $\mathcal{O}(1/\sqrt{T})$ without any dependence on the dimensions. Such improved convergence rate is under the assumption that the gradient estimator have a bounded L_2 norm, which can thus hide a dependency of $\sqrt{d}$ in the final convergence rate.

3.3 Convergence Analysis

Proof of Theorem 2. The detailed proof of this section is reported in the supplementary material. We now present a proof sketch for out main convergence result of Algorithm 2.

Step 1: Reparameterization. Similarly to [Yan et al., 2018, Chen et al., 2018] with SGD (with momentum) and centralized adaptive gradient methods, define the following auxiliary sequence:

$$Z_t = \overline{X}_t + \frac{\beta_1}{1-\beta_1}(\overline{X}_t - \overline{X}_{t-1}), \tag{3}$$

with $\overline{X}_0 \triangleq \overline{X}_1$. Such an auxiliary sequence can help us deal with the bias brought by the momentum and simplifies the convergence analysis. An intermediary result needed to conduct our proof reads:

Lemma 1. *For the sequence defined in (3), we have*

$$Z_{t+1} - Z_t = \alpha \frac{\beta_1}{1-\beta_1} \frac{1}{N} \sum_{i=1}^{N} m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}}\right) - \alpha \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \tag{4}$$

Lemma 1 does not display any momentum term in $\frac{1}{N}\sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}}$. This simplification is convenient since it is directly related to the current gradients instead of the exponential average of past gradients.

Step 2: Smoothness. Using smoothness assumption A1 involves the following scalar product term: $\kappa_t := \langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N} \nabla f_i(x_{t,i})/\sqrt{\overline{U}_t} \rangle$ which can be lower bounded by:

$$\kappa_t \geq \frac{1}{2} \left\| \frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 - \frac{3}{2} \left\| \frac{\nabla f(Z_t) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 - \frac{3}{2} \left\| \frac{\frac{1}{N}\sum_{i=1}^{N} \nabla f_i(x_{t,i}) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 . \tag{5}$$

The above inequality substituted in the smoothness condition $f(Z_{t+1}) \leq f(Z_t) + \langle \nabla f(Z_t), Z_{t+1} - Z_t \rangle + \frac{L}{2}\|Z_{t+1} - Z_t\|^2$ yields:

$$\frac{1}{T}\sum_{t=1}^{T} \mathbb{E}\left[\left\| \frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 \right] \leq \frac{2}{T\alpha}(\mathbb{E}[\Delta_f]) + \frac{L}{T\alpha} \sum_{t=1}^{T} \mathbb{E}\left[\|Z_{t+1} - Z_t\|^2\right] + \frac{2}{T}\frac{\beta_1}{1-\beta_1} T_1 + \frac{2}{T} T_2 + \frac{3}{T} T_3 \,, \tag{6}$$

where $\Delta_f := \mathbb{E}[f(Z_1)] - \mathbb{E}[f(Z_{T+1})]$ T_1, T_2 and T_3 are three terms, defined in the supplementary material, and which can be tightly bounded from above. We first bound T_3 using the following quantities of interest:

$$\sum_{t=1}^{T} \left\|Z_t - \overline{X}_t\right\|^2 \leq T \left(\frac{\beta_1}{1-\beta_1}\right)^2 \alpha^2 d \frac{G_\infty^2}{\epsilon} \quad \text{and} \quad \sum_{t=1}^{T} \frac{1}{N} \sum_{i=1}^{N} \left\|x_{t,i} - \overline{X}_t\right\|^2 \leq T\alpha^2 \left(\frac{1}{1-\lambda}\right)^2 d G_\infty^2 \frac{1}{\epsilon} \,. \tag{7}$$

where $\lambda = \max(|\lambda_2|, |\lambda_N|)$ and recall that λ_i is ith largest eigenvalue of W.

Then, concerning the term T_2, few derivations, not detailed here for simplicity, yields:

$$T_2 \leq \frac{G_\infty^2}{N} \mathbb{E}\left[\sum_{t=1}^{T} \frac{1}{2\epsilon^{1.5}} \| - \sum_{l=2}^{N} \tilde{U}_t q_l q_l^T \|_{abs} \right] \tag{8}$$

where q_l is the eigenvector corresponding to lth largest eigenvalue of W and $\|\cdot\|_{abs}$ is the entry-wise L_1 norm of matrices. We can also show that

$$\sum_{t=1}^{T} \| - \sum_{l=2}^{N} \tilde{U}_t q_l q_l^T \|_{abs} \leq \sqrt{N} \sum_{o=0}^{T-1} \frac{\lambda}{1-\lambda} \|(-\hat{V}_{o-1} + \hat{V}_o)\|_{abs} \tag{9}$$

resulting in an upper bound for T_2 proportional to $\sum_{o=0}^{T-1} \|(-\hat{V}_{o-1} + \hat{V}_o)\|_{abs}$. Similarly:

$$T_1 \leq G_\infty^2 \frac{1}{2\epsilon^{1.5}} \frac{1}{\sqrt{N}} \mathbb{E}\left[\frac{1}{1-\lambda} \sum_{t=1}^{T} \|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs} \right] \tag{10}$$

Step 3: Bounding the drift term variance. An important term that needs upper bounding in our proof is the variance of the gradients multiplied (element-wise) by the adaptive learning rate:

$$\mathbb{E}\left[\left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \right] \leq \mathbb{E}[\|\Gamma_u^f\|^2] + \frac{d}{N}\frac{\sigma^2}{\epsilon} \tag{11}$$

where $\Gamma_u^f := 1/N \sum_{i=1}^{N} \nabla f_i(x_{t,i})/\sqrt{u_{t,i}}$. Two consecutive and simple bounding of the above yields:

$$\sum_{t=1}^{T} \mathbb{E}[\|\Gamma_u^f\|^2] \leq 2 \sum_{t=1}^{T} \mathbb{E}[\|\Gamma_{\overline{U}}^f\|^2] + 2 \sum_{t=1}^{T} \mathbb{E}\left[\frac{1}{N} \sum_{i=1}^{N} G_\infty^2 \frac{1}{\sqrt{\epsilon}} \left\| \frac{1}{\sqrt{u_{t,i}}} - \frac{1}{\sqrt{\overline{U}_t}} \right\|_1 \right] \tag{12}$$

and

$$\sum_{t=1}^{T} \mathbb{E}[\|\Gamma_{\overline{U}}^f\|^2] \leq 2 \sum_{t=1}^{T} \mathbb{E}\left[\left\| \frac{\nabla f(\overline{X}_t)}{\sqrt{\overline{U}_t}} \right\|^2 \right] + 2 \sum_{t=1}^{T} \mathbb{E}\left[\left\| \frac{1}{N} \sum_{i=1}^{N} \frac{\nabla f_i(\overline{X}_t) - \nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}} \right\|^2 \right] . \tag{13}$$

Then, by plugging the LHS of (13) in (6), and further bounding as operated for T_2, T_3 (see supplement), we obtain the bound in Theorem 2.

4 Numerical Experiments

In this section, we conduct experiments to test the performance of Decentralized AMSGrad, see Algorithm 3, on both homogeneous data and heterogeneous data distribution (i.e. the data generating distribution on different nodes are assumed to be different). We compare it with DADAM and the decentralized stochastic gradient descent (DGD) developed in [Lian et al., 2017]. The task consists of training a Convolutional Neural Network (CNN) with 3 convolution layers followed by a fully connected layer on MNIST [LeCun, 1998]. We set $\epsilon = 10^{-6}$ for both Decentralized AMSGrad and DADAM, the learning rate is chosen from the grid $[10^{-1}, 10^{-2}, 10^{-3}, 10^{-4}, 10^{-5}, 10^{-6}]$ based on validation accuracy for all algorithms. In all the following experiments, the graph contains 5 nodes and the nodes form a ring, each node can only talk with its two adjacent neighbors. We set $W_{ij} = 1/3$ if there nodes i and j are neighbors and $W_{ij} = 0$ otherwise for the mixing matrix. More details and experiments can be found in Appendix A.4.

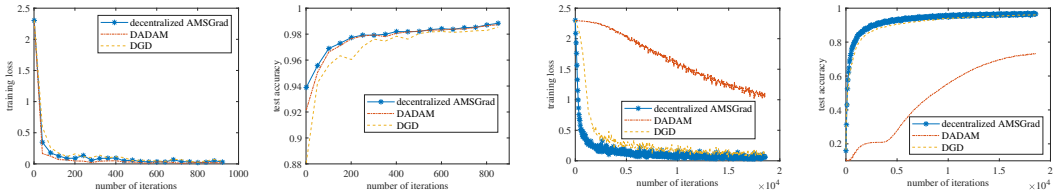

Figure 1: Performance comparison on homogeneous and heterogeneous data

Figure 1 shows the performance of different algorithms on homogeneous data. The whole dataset is shuffled evenly split to different nodes. We can see that decentralized AMSGrad and DADAM performs quite similarly and DGD is slower compared with them in terms of both training loss and test accuracy. Though we have prove in previous sections that DADAM is not a convergent algorithm, its performance is still quite good on homogeneous data. The reason is that the adaptive learning rates tend to be similar on different nodes when we have homogeneous data distribution. However, this is usually not true when we have heterogeneous data distribution. This motivates us to compare the performance of the algorithms on a different data distribution.

In Figure 1, we compare the performance of different algorithms on heterogeneous data. In this case, each node only contains training data with two labels out of ten. We can see that all algorithm converges significantly slower compared with the case with homogeneous data. Especially, the performance of DADAM deteriorates significantly. decentralized AMSGrad achieves the best training and testing performance in this experiment.

5 Conclusion

This paper studies the problem of designing adaptive gradient methods for decentralized training. We propose a unifying algorithmic framework that can convert existing adaptive gradient methods to decentralized settings. With rigorous convergence analysis, we show that if the original algorithm satisfies converges under some minor conditions, the converted algorithm obtained using our proposed framework is guaranteed to converge to stationary points of the regret function. By applying our framework to AMSGrad, we propose the first convergent adaptive gradient methods, namely Decentralized AMSGrad. Experiments show that the proposed algorithm achieves better performance than the baselines.

6 Broader Impact of Our Work

References

Naman Agarwal, Brian Bullins, Xinyi Chen, Elad Hazan, Karan Singh, Cyril Zhang, and Yi Zhang. The case for full-matrix adaptive regularization. *arXiv preprint arXiv:1806.02958*, 2018.

Alham Fikri Aji and Kenneth Heafield. Sparse communication for distributed gradient descent. *arXiv preprint arXiv:1704.05021*, 2017.

Dan Alistarh, Demjan Grubic, Jerry Li, Ryota Tomioka, and Milan Vojnovic. Qsgd: Communication-efficient sgd via gradient quantization and encoding. In *Advances in Neural Information Processing Systems*, pages 1709–1720, 2017.

Mahmoud Assran, Nicolas Loizou, Nicolas Ballas, and Michael Rabbat. Stochastic gradient push for distributed deep learning. *arXiv preprint arXiv:1811.10792*, 2018.

Stephen Boyd, Neal Parikh, Eric Chu, Borja Peleato, Jonathan Eckstein, et al. Distributed optimization and statistical learning via the alternating direction method of multipliers. *Foundations and Trends® in Machine learning*, 3(1):1–122, 2011.

Xiangyi Chen, Sijia Liu, Ruoyu Sun, and Mingyi Hong. On the convergence of a class of adam-type algorithms for non-convex optimization. *arXiv preprint arXiv:1808.02941*, 2018.

Yongjian Chen, Tao Guan, and Cheng Wang. Approximate nearest neighbor search by residual vector quantization. *Sensors*, 10(12):11259–11273, 2010.

Trishul Chilimbi, Yutaka Suzue, Johnson Apacible, and Karthik Kalyanaraman. Project adam: Building an efficient and scalable deep learning training system. In *11th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 14)*, pages 571–582, 2014.

Paolo Di Lorenzo and Gesualdo Scutari. Next: In-network nonconvex optimization. *IEEE Transactions on Signal and Information Processing over Networks*, 2(2):120–136, 2016.

John Duchi, Elad Hazan, and Yoram Singer. Adaptive subgradient methods for online learning and stochastic optimization. *Journal of Machine Learning Research*, 12(Jul):2121–2159, 2011a.

John C Duchi, Alekh Agarwal, and Martin J Wainwright. Dual averaging for distributed optimization: Convergence analysis and network scaling. *IEEE Transactions on Automatic control*, 57(3): 592–606, 2011b.

Tiezheng Ge, Kaiming He, Qifa Ke, and Jian Sun. Optimized product quantization for approximate nearest neighbor search. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2946–2953, 2013.

Mingyi Hong, Davood Hajinezhad, and Ming-Min Zhao. Prox-pda: The proximal primal-dual algorithm for fast distributed nonconvex optimization and learning over networks. In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pages 1529–1538. JMLR. org, 2017.

Herve Jegou, Matthijs Douze, and Cordelia Schmid. Product quantization for nearest neighbor search. *IEEE transactions on pattern analysis and machine intelligence*, 33(1):117–128, 2010.

Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.

Yann LeCun. The mnist database of handwritten digits. *http://yann. lecun. com/exdb/mnist/*, 1998.

Xiaoyu Li and Francesco Orabona. On the convergence of stochastic gradient descent with adaptive stepsizes. *arXiv preprint arXiv:1805.08114*, 2018.

Xiangru Lian, Ce Zhang, Huan Zhang, Cho-Jui Hsieh, Wei Zhang, and Ji Liu. Can decentralized algorithms outperform centralized algorithms? a case study for decentralized parallel stochastic gradient descent. In *Advances in Neural Information Processing Systems*, pages 5330–5340, 2017.

Yujun Lin, Song Han, Huizi Mao, Yu Wang, and William J Dally. Deep gradient compression: Reducing the communication bandwidth for distributed training. *arXiv preprint arXiv:1712.01887*, 2017.

Liangchen Luo, Yuanhao Xiong, Yan Liu, and Xu Sun. Adaptive gradient methods with dynamic bound of learning rate. *arXiv preprint arXiv:1902.09843*, 2019.

H Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, et al. Communication-efficient learning of deep networks from decentralized data. *arXiv preprint arXiv:1602.05629*, 2016.

Parvin Nazari, Davoud Ataee Tarzanagh, and George Michailidis. Dadam: A consensus-based distributed adaptive gradient method for online optimization. *arXiv preprint arXiv:1901.09109*, 2019.

Angelia Nedic and Asuman Ozdaglar. Distributed subgradient methods for multi-agent optimization. *IEEE Transactions on Automatic Control*, 54(1):48, 2009.

Sashank Reddi, Zachary Charles, Manzil Zaheer, Zachary Garrett, Keith Rush, Jakub Konečnỳ, Sanjiv Kumar, and H Brendan McMahan. Adaptive federated optimization. *arXiv preprint arXiv:2003.00295*, 2020.

Sashank J Reddi, Satyen Kale, and Sanjiv Kumar. On the convergence of adam and beyond. *arXiv preprint arXiv:1904.09237*, 2019.

Wei Shi, Qing Ling, Gang Wu, and Wotao Yin. Extra: An exact first-order algorithm for decentralized consensus optimization. *SIAM Journal on Optimization*, 25(2):944–966, 2015.

Sebastian U Stich, Jean-Baptiste Cordonnier, and Martin Jaggi. Sparsified sgd with memory. In *Advances in Neural Information Processing Systems*, pages 4447–4458, 2018.

Hanlin Tang, Xiangru Lian, Ming Yan, Ce Zhang, and Ji Liu. D^2: Decentralized training over decentralized data. *arXiv preprint arXiv:1803.07068*, 2018.

Hanlin Tang, Xiangru Lian, Tong Zhang, and Ji Liu. Doublesqueeze: Parallel stochastic gradient descent with double-pass error-compensated compression. *arXiv preprint arXiv:1905.05957*, 2019.

Hongyi Wang, Scott Sievert, Shengchao Liu, Zachary Charles, Dimitris Papailiopoulos, and Stephen Wright. Atomo: Communication-efficient learning via atomic sparsification. In *Advances in Neural Information Processing Systems*, pages 9850–9861, 2018.

Jianqiao Wangni, Jialei Wang, Ji Liu, and Tong Zhang. Gradient sparsification for communication-efficient distributed optimization. In *Advances in Neural Information Processing Systems*, pages 1299–1309, 2018.

Rachel Ward, Xiaoxia Wu, and Leon Bottou. Adagrad stepsizes: Sharp convergence over nonconvex landscapes, from any initialization. *arXiv preprint arXiv:1806.01811*, 2018.

Yan Yan, Tianbao Yang, Zhe Li, Qihang Lin, and Yi Yang. A unified analysis of stochastic momentum methods for deep learning. *arXiv preprint arXiv:1808.10396*, 2018.

Kun Yuan, Qing Ling, and Wotao Yin. On the convergence of decentralized gradient descent. *SIAM Journal on Optimization*, 26(3):1835–1854, 2016.

Manzil Zaheer, Sashank Reddi, Devendra Sachan, Satyen Kale, and Sanjiv Kumar. Adaptive methods for nonconvex optimization. In *Advances in Neural Information Processing Systems*, pages 9793–9803, 2018.

Fangyu Zou and Li Shen. On the convergence of weighted adagrad with momentum for training deep neural networks. *arXiv preprint arXiv:1808.03408*, 2018.

A Appendix

A.1 Proof of Theorem 2

To prove convergence of the algorithm, we first define an auxiliary sequence

$$Z_t = \overline{X}_t + \frac{\beta_1}{1-\beta_1}(\overline{X}_t - \overline{X}_{t-1}) \tag{14}$$

with $\overline{X}_0 \triangleq \overline{X}_1$.

Then we have the following Lemma to characterize the difference of iterations of sequence Z_t.

Lemma. *For the sequence defined in (14), we have*

$$Z_{t+1} - Z_t = \alpha\frac{\beta_1}{1-\beta_1}\frac{1}{N}\sum_{i=1}^{N} m_{t-1,i} \odot (\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}}) - \alpha\frac{1}{N}\sum_{i=1}^{N}\frac{g_{t,i}}{\sqrt{u_{t,i}}} \tag{15}$$

Proof: See Appendix A.3. □

Since $\mathbb{E}[g_{t,i}] = \nabla f(x_{t,i})$ and $u_{t,i}$ is a function of $G_{1:t-1}$ (which denotes $G_1, G_2, \cdots, G_{t-1}$), we have

$$\mathbb{E}_{G_t|G_{1:t-1}}\left[\frac{1}{N}\sum_{i=1}^{N}\frac{g_{t,i}}{\sqrt{u_{t,i}}}\right] = \frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}} \tag{16}$$

By assuming smoothness (A1) we have

$$f(Z_{t+1}) \le f(Z_t) + \langle \nabla f(Z_t), Z_{t+1} - Z_t\rangle + \frac{L}{2}\|Z_{t+1} - Z_t\|^2 \tag{17}$$

Substitute (61) into the above inequality and take expectation over G_t given $G_{1:t-1}$, we have

$$\begin{aligned}\mathbb{E}_{G_t|G_{1:t-1}}[f(Z_{t+1})] \le & f(Z_t) - \alpha\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\rangle + \frac{L}{2}\mathbb{E}_{G_t|G_{1:t-1}}\left[\|Z_{t+1} - Z_t\|^2\right] \\ & + \alpha\frac{\beta_1}{1-\beta_1}\mathbb{E}_{G_t|G_{1:t-1}}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N} m_{t-1,i}\odot(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}})\right\rangle\right]\end{aligned} \tag{18}$$

Then take expectation over $G_{1:t-1}$ and rearrange, we have

$$\begin{aligned}\alpha\mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\rangle\right] \le & \mathbb{E}[f(Z_t)] - \mathbb{E}[f(Z_{t+1})] + \frac{L}{2}\mathbb{E}\left[\|Z_{t+1} - Z_t\|^2\right] \\ & + \alpha\frac{\beta_1}{1-\beta_1}\mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N} m_{t-1,i}\odot(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}})\right\rangle\right]\end{aligned} \tag{19}$$

In addition, we have

$$\begin{aligned}&\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\rangle \\ =&\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\rangle + \left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N}\nabla f_i(x_{t,i})\odot\left(\frac{1}{\sqrt{u_{t,i}}} - \frac{1}{\sqrt{\overline{U}_t}}\right)\right\rangle\end{aligned} \tag{20}$$

and the first term on RHS of the equality can be lower bounded as

$$
\begin{aligned}
&\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}} \right\rangle \\
=&\frac{1}{2}\left\| \frac{\nabla f(Z_t)}{\overline{U}_t^{1/4}} \right\|^2 + \frac{1}{2}\left\| \frac{\frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i})}{\overline{U}_t^{1/4}} \right\|^2 - \frac{1}{2}\left\| \frac{\nabla f(Z_t) - \frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i})}{\overline{U}_t^{1/4}} \right\|^2 \\
\geq&\frac{1}{4}\left\| \frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 + \frac{1}{4}\left\| \frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 - \frac{1}{2}\left\| \frac{\nabla f(Z_t) - \frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i})}{\overline{U}_t^{1/4}} \right\|^2 \\
&- \frac{1}{2}\left\| \frac{\nabla f(Z_t) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 - \frac{1}{2}\left\| \frac{\frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i}) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 \\
\geq&\frac{1}{2}\left\| \frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 - \frac{3}{2}\left\| \frac{\nabla f(Z_t) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 - \frac{3}{2}\left\| \frac{\frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i}) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2
\end{aligned} \tag{21}
$$

where the inequalities are all due to Cauchy-Schwartz.

Substituting (21) and (20) into (19), we get

$$
\begin{aligned}
\frac{1}{2}\alpha\mathbb{E}\left[\left\| \frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2\right] \leq& \mathbb{E}[f(Z_t)] - \mathbb{E}[f(Z_{t+1})] + \frac{L}{2}\mathbb{E}\left[\|Z_{t+1} - Z_t\|^2\right] \\
&+ \alpha\frac{\beta_1}{1-\beta_1}\mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^N m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}}\right)\right\rangle\right] \\
&- \alpha\mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i}) \odot \left(\frac{1}{\sqrt{u_{t,i}}} - \frac{1}{\sqrt{\overline{U}_t}}\right)\right\rangle\right] \\
&+ \frac{3}{2}\alpha\mathbb{E}\left[\left\| \frac{\frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i}) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 + \left\| \frac{\nabla f(Z_t) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2\right]
\end{aligned} \tag{22}
$$

Then sum over the above inequality from $t = 1$ to T and divide both sides by $T\alpha/2$, we have

$$
\begin{aligned}
\frac{1}{T}\sum_{t=1}^T \mathbb{E}\left[\left\| \frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2\right] \leq& \frac{2}{T\alpha}(\mathbb{E}[f(Z_1)] - \mathbb{E}[f(Z_{T+1})]) + \frac{L}{T\alpha}\sum_{t=1}^T \mathbb{E}\left[\|Z_{t+1} - Z_t\|^2\right] \\
&+ \frac{2}{T}\frac{\beta_1}{1-\beta_1}\underbrace{\sum_{t=1}^T \mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^N m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}}\right)\right\rangle\right]}_{T_1} \\
&+ \frac{2}{T}\underbrace{\sum_{t=1}^T \mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i}) \odot \left(\frac{1}{\sqrt{\overline{U}_t}} - \frac{1}{\sqrt{u_{t,i}}}\right)\right\rangle\right]}_{T_2} \\
&+ \frac{3}{T}\underbrace{\sum_{t=1}^T \mathbb{E}\left[\left\| \frac{\frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i}) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2 + \left\| \frac{\nabla f(Z_t) - \nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}} \right\|^2\right]}_{T_3}
\end{aligned} \tag{23}
$$

Now we need to upper bound all the terms on RHS of the above inequality to get the convergence rate.

For terms in T_3 in (23), we can upper bound them by

$$\left\|\frac{\nabla f(Z_t)-\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2 \leq \frac{1}{\min_{j\in[d]}[\overline{U}_t^{1/2}]_j}\left\|\nabla f(Z_t)-\nabla f(\overline{X}_t)\right\|^2 \leq L\frac{1}{\min_{j\in[d]}[\overline{U}_t^{1/2}]_j}\underbrace{\left\|Z_t-\overline{X}_t\right\|^2}_{T_4} \tag{24}$$

and

$$\begin{aligned}\left\|\frac{\frac{1}{N}\sum_{i=1}^N\nabla f_i(x_{t,i})-\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2 \leq& \frac{1}{\min_{j\in[d]}[\overline{U}_t^{1/2}]_j}\frac{1}{N}\sum_{i=1}^N\left\|\nabla f_i(x_{t,i})-\nabla f(\overline{X}_t)\right\|^2 \\ \leq& L\frac{1}{\min_{j\in[d]}[\overline{U}_t^{1/2}]_j}\frac{1}{N}\underbrace{\sum_{i=1}^N\left\|x_{t,i}-\overline{X}_t\right\|^2}_{T_5}\end{aligned} \tag{25}$$

using Jensen's inequality, Lipschitz continuity of f_i, and the fact that $f=\frac{1}{N}\sum_{i=1}^N f_i$. .

What we need to do next is to bound T_4 and T_5 and we will bound T_5 first.

Before we proceed into bounding T_5, we need some preparations. Let's recall the update rule of X_t, we have

$$X_t = X_{t-1}W-\alpha\frac{M_{t-1}}{\sqrt{U_{t-1}}} = X_1W^{t-1}-\alpha\sum_{k=0}^{t-2}\frac{M_{t-k-1}}{\sqrt{U_{t-k-1}}}W^k \tag{26}$$

where we define $W^0=\mathbf{I}$.

Since W is a symmetric matrix, we can decompose it as $W=Q\Lambda Q^T$ where Q is a orthonormal matrix and Λ is a diagonal matrix whose diagonal elements correspond to eigenvalues of W in an descending order, i.e. $\Lambda_{ii}=\lambda_i$ with λ_i being ith largest eigenvalue of W. In addition, because W is a doubly stochastic matrix, we know $\lambda_1=1$ and $q_1=\frac{\mathbf{1}_N}{\sqrt{N}}$

With eigen-decomposition of W, we can rewrite T_5 as

$$\sum_{i=1}^N\left\|x_{t,i}-\overline{X}_t\right\|^2 = \|X_t-\overline{X}_t\mathbf{1}_N^T\|_F^2 = \|X_tQQ^T-X_t\frac{1}{N}\mathbf{1}_N\mathbf{1}_N^T\|_F^2 = \sum_{l=2}^N\|X_tq_l\|^2 \tag{27}$$

In addition, we can rewrite (26) as

$$X_t = X_1W^{t-1}-\alpha\sum_{k=0}^{t-2}\frac{M_{t-k-1}}{\sqrt{U_{t-k-1}}}W^k = X_1-\alpha\sum_{k=0}^{t-2}\frac{M_{t-k-1}}{\sqrt{U_{t-k-1}}}Q\Lambda^kQ^T \tag{28}$$

where the last equality is because $x_{1,i}=x_{1,j},\forall i,j$ and thus $X_1W=X_1$.

Then we have when $l>1$,

$$X_tq_l = (X_1-\alpha\sum_{k=0}^{t-2}\frac{M_{t-k-1}}{\sqrt{U_{t-k-1}}}Q\Lambda^kQ^T)q_l = -\alpha\sum_{k=0}^{t-2}\frac{M_{t-k-1}}{\sqrt{U_{t-k-1}}}q_l\lambda_l^k \tag{29}$$

because Q is orthonormal and $X_1q_l=x_{1,1}\mathbf{1}_N^Tq_l=x_{1,1}\sqrt{N}q_1^Tq_l=0,\forall l\neq 1$.

Combining (27) and (29), we have

$$T_5=\sum_{i=1}^N\left\|x_{t,i}-\overline{X}_t\right\|^2=\sum_{l=2}^N\|X_tq_l\|^2=\sum_{l=2}^N\alpha^2\left\|\sum_{k=0}^{t-2}\frac{M_{t-k-1}}{\sqrt{U_{t-k-1}}}\lambda_l^kq_l\right\|^2 \leq \alpha^2\left(\frac{1}{1-\lambda}\right)^2NdG_\infty^2\frac{1}{\epsilon} \tag{30}$$

where the last inequality follows from the fact that $g_{t,i}\leq G_\infty$, $\|q_l\|=1$, and $|\lambda_l|\leq\lambda<1$.

Now let us turn to T_4, it can be rewritten as

$$\left\|Z_t - \overline{X}_t\right\|^2 = \left\|\frac{\beta_1}{1-\beta_1}(\overline{X}_t - \overline{X}_{t-1})\right\|^2 = \left(\frac{\beta_1}{1-\beta_1}\right)^2 \alpha^2 \left\|\frac{1}{N}\sum_{i=1}^{N}\frac{m_{t-1,i}}{\sqrt{u_{t-1,i}}}\right\|^2 \leq \left(\frac{\beta_1}{1-\beta_1}\right)^2 \alpha^2 d\frac{G_\infty^2}{\epsilon} \tag{31}$$

Now we know both T_4 and T_5 are in the order of $\mathcal{O}\alpha^2)$ and thus T_3 is in the order of $\mathcal{O}\alpha^2)$.

Next we will bound T_2 and T_1. Define $G_1 \triangleq \max_{t\in[T]}\max_{i\in[N]}\|\nabla f_i(x_{t,i})\|_\infty$, $G_2 \triangleq \max_{t\in[T]}\|\nabla f(Z_t)\|_\infty$, $G_3 \triangleq \max_{t\in[T]}\max_{i\in[N]}\|g_{t,i}\|_\infty$ and $G_\infty = \max(G_1, G_2, G_3)$

Then we have

$$\begin{aligned}
T_2 =& \sum_{t=1}^{T}\mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N}\nabla f_i(x_{t,i}) \odot \left(\frac{1}{\sqrt{\overline{U}_t}} - \frac{1}{\sqrt{u_{t,i}}}\right)\right\rangle\right] \\
\leq& \sum_{t=1}^{T}\mathbb{E}\left[G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\left|\frac{1}{\sqrt{[\overline{U}_t]_j}} - \frac{1}{\sqrt{[u_{t,i}]_j}}\right|\right] \\
=& \sum_{t=1}^{T}\mathbb{E}\left[G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\left|\frac{1}{\sqrt{[\overline{U}_t]_j}} - \frac{1}{\sqrt{[u_{t,i}]_j}}\right|\frac{\sqrt{[\overline{U}_t]_j} + \sqrt{[u_{t,i}]_j}}{\sqrt{[\overline{U}_t]_j} + \sqrt{[u_{t,i}]_j}}\right] \\
=& \sum_{t=1}^{T}\mathbb{E}\left[G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\left|\frac{[\overline{U}_t]_j - [u_{t,i}]_j}{[\overline{U}_t]_j\sqrt{[u_{t,i}]_j} + \sqrt{[\overline{U}_t]_j}[u_{t,i}]_j}\right|\right] \\
\leq& \mathbb{E}\Bigg[\underbrace{\sum_{t=1}^{T} G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\left|\frac{[\overline{U}_t]_j - [u_{t,i}]_j}{2\epsilon^{1.5}}\right|}_{T_6}\Bigg]
\end{aligned} \tag{32}$$

where the last inequality is due to $[u_{t,i}]_j \geq \epsilon,\ \forall t, i, j$.

To simplify notations, let's define $\|A\|_{abs} = \sum_{i,j}|A_{ij}|$ to be the entry-wise L_1 norm of a matrix A, then we have

$$\begin{aligned}
T_6 \leq& \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\|\overline{U}_t\mathbf{1}^T - U_t\|_{abs} \\
\leq& \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\|\overline{\tilde{U}}_t\mathbf{1}^T - \tilde{U}_t\|_{abs} \\
=& \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\|\tilde{U}_t\frac{1}{N}\mathbf{1}_N\mathbf{1}_N^T - \tilde{U}_t Q Q^T\|_{abs} \\
=& \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\|-\tilde{U}_t\sum_{l=2}^{N} q_l q_l^T\|_{abs} \\
=& \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\|-\sum_{l=2}^{N}\tilde{U}_t q_l q_l^T\|_{abs}
\end{aligned}$$

where the second inequality is due to Lemma 2 and the fact that $U_t = \max(\tilde{U}_t, \epsilon)$ element-wisely.

Lemma 2. *Given a set of numbers $a_1, \cdots, a_n$ and denote their mean to be $\bar{a} = \frac{1}{n}\sum_{i=1}^{n} a_i$. In addition, define $b_i(r) \triangleq= \max(a_i, r)$ and $\bar{b}(r) = \frac{1}{n}\sum_{i=1}^{n} b_i(r)$. For any r and r' with $r' \geq r$ we have*

$$\sum_{i=1}^{n}|b_i(r) - \bar{b}(r)| \geq \sum_{i=1}^{n}|b_i(r') - \bar{b}(r')| \tag{33}$$

and when $r \leq \min_{i\in[n]} a_i$, we have

$$\sum_{i=1}^{n} |b_i(r) - \bar{b}(r)| = \sum_{i=1}^{n} |a_i - \bar{a}| \tag{34}$$

Proof: See Appendix A.3. □

Recall from update rule of U_t, by defining $\hat{V}_{-1} \triangleq \hat{V}_0$ and $U_0 \triangleq U_{1/2}$, we have $\forall t \geq 0$

$$\tilde{U}_{t+1} = (\tilde{U}_t - \hat{V}_{t-1} + \hat{V}_t)W \tag{35}$$

and thus

$$\tilde{U}_t = \tilde{U}_0 W^t + \sum_{k=1}^{t}(-\hat{V}_{t-1-k} + \hat{V}_{t-k})W^k = \tilde{U}_0 + \sum_{k=1}^{t}(-\hat{V}_{t-1-k} + \hat{V}_{t-k})Q\Lambda^k Q^T \tag{36}$$

Then we further have when $l \neq 1$,

$$\tilde{U}_t q_l = (\tilde{U}_0 + \sum_{k=1}^{t}(-\hat{V}_{t-1-k} + \hat{V}_{t-k})Q\Lambda^k Q^T)q_l = \sum_{k=1}^{t}(-\hat{V}_{t-1-k} + \hat{V}_{t-k})q_l\lambda_l^k \tag{37}$$

where the last equality is due to the definition $\tilde{U}_0 \triangleq U_{1/2} = \epsilon \mathbf{1_d}\mathbf{1}_N^T = \sqrt{N}\epsilon \mathbf{1_d}\mathbf{1}_N^T$ (recall that $q_1 = \frac{1}{\sqrt{N}}\mathbf{1}_N^T$) and $q_i^T q_j = 0$ when $i \neq j$.

Note by definition of $\|\cdot\|_{abs}$, we have $\forall A, B, \|A+B\|_{abs} \leq \|A\|_{abs} + \|B\|_{abs}$, then we have

$$\begin{aligned}
T_6 \leq & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\| - \sum_{l=2}^{N}\tilde{U}_t q_l q_l^T\|_{abs} \\
= & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\| - \sum_{k=1}^{t}(-\hat{V}_{t-1-k} + \hat{V}_{t-k})\sum_{l=2}^{N} q_l\lambda_l^k q_l^T\|_{abs} \\
\leq & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\sum_{k=1}^{t}\|(-\hat{V}_{t-1-k} + \hat{V}_{t-k})\sum_{l=2}^{N} q_l\lambda_l^k q_l^T\|_{abs} \\
= & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\sum_{k=1}^{t}\sum_{j=1}^{d}\|\sum_{l=2}^{N} q_l\lambda_l^k q_l^T(-\hat{V}_{t-1-k} + \hat{V}_{t-k})^T e_j\|_1 \\
\leq & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\sum_{k=1}^{t}\sum_{j=1}^{d}\|\sum_{l=2}^{N} q_l\lambda_l^k q_l^T\|_1\|(-\hat{V}_{t-1-k} + \hat{V}_{t-k})^T e_j\|_1 \\
\leq & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\sum_{k=1}^{t}\sum_{j=1}^{d}\sqrt{N}\|\sum_{l=2}^{N} q_l\lambda_l^k q_l^T\|_2\|(-\hat{V}_{t-1-k} + \hat{V}_{t-k})^T e_j\|_1 \\
\leq & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\sum_{k=1}^{t}\sum_{j=1}^{d}\|(-\hat{V}_{t-1-k} + \hat{V}_{t-k})^T e_j\|_1\sqrt{N}\lambda^k \\
= & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\sum_{k=1}^{t}\|(-\hat{V}_{t-1-k} + \hat{V}_{t-k})\|_{abs}\sqrt{N}\lambda^k \\
= & \frac{G_\infty^2}{N}\sum_{t=1}^{T}\frac{1}{2\epsilon^{1.5}}\sum_{o=0}^{t-1}\|(-\hat{V}_{o-1} + \hat{V}_o)\|_{abs}\sqrt{N}\lambda^{t-o} \\
= & \frac{G_\infty^2}{N}\frac{1}{2\epsilon^{1.5}}\sum_{o=0}^{T-1}\sum_{t=o+1}^{T}\|(-\hat{V}_{o-1} + \hat{V}_o)\|_{abs}\sqrt{N}\lambda^{t-o} \\
\leq & \frac{G_\infty^2}{\sqrt{N}}\frac{1}{2\epsilon^{1.5}}\sum_{o=0}^{T-1}\frac{\lambda}{1-\lambda}\|(-\hat{V}_{o-1} + \hat{V}_o)\|_{abs}
\end{aligned} \tag{38}$$

where $\lambda = \max(|\lambda_2|, |\lambda_N|)$.

Combining (32) and (38), we have

$$T_2 \leq \frac{G_\infty^2}{\sqrt{N}} \frac{1}{2\epsilon^{1.5}} \frac{\lambda}{1-\lambda} \mathbb{E}\left[\sum_{o=0}^{T-1} \|(-\hat{V}_{o-1} + \hat{V}_o)\|_{abs}\right] \tag{39}$$

Now we need to bound T_1, we have

$$\begin{aligned}
T_1 =& \sum_{t=1}^{T} \mathbb{E}\left[\left\langle \nabla f(Z_t), \frac{1}{N}\sum_{i=1}^{N} m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}}\right)\right\rangle\right] \\
\leq& \sum_{t=1}^{T} \mathbb{E}\left[G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\left|\frac{1}{\sqrt{[u_{t-1,i}]_j}} - \frac{1}{\sqrt{[u_{t,i}]_j}}\right|\right] \\
=& \sum_{t=1}^{T} \mathbb{E}\left[G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\left|\left(\frac{1}{\sqrt{[u_{t-1,i}]_j}} - \frac{1}{\sqrt{[u_{t,i}]_j}}\right)\frac{\sqrt{[u_{t,i}]_j} + \sqrt{[u_{t-1,i}]_j}}{\sqrt{[u_{t,i}]_j} + \sqrt{[u_{t-1,i}]_j}}\right|\right] \\
\leq& \sum_{t=1}^{T} \mathbb{E}\left[G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\left|\frac{1}{2\epsilon^{1.5}}([u_{t-1,i}]_j - [u_{t,i}]_j)\right|\right] \\
\overset{(a)}{\leq}& \sum_{t=1}^{T} \mathbb{E}\left[G_\infty^2 \frac{1}{N}\sum_{i=1}^{N}\sum_{j=1}^{d}\frac{1}{2\epsilon^{1.5}}|([\tilde{u}_{t-1,i}]_j - [\tilde{u}_{t,i}]_j)|\right] \\
=& G_\infty^2 \frac{1}{2\epsilon^{1.5}} \frac{1}{N} \mathbb{E}\left[\sum_{t=1}^{T} \|\tilde{U}_{t-1} - \tilde{U}_t\|_{abs}\right]
\end{aligned} \tag{40}$$

where (a) is due to $[\tilde{u}_{t-1,i}]_j = \max([u_{t-1,i}]_j, \epsilon)$ and the function $\max(\cdot, \epsilon)$ is 1-Lipschitz.

In addition, by update rule of U_t, we have

$$\begin{aligned}
&\sum_{t=1}^{T} \|\tilde{U}_{t-1} - \tilde{U}_t\|_{abs} \\
=& \sum_{t=1}^{T} \|\tilde{U}_{t-1} - (\tilde{U}_{t-1} - \hat{V}_{t-2} + \hat{V}_{t-1})W\|_{abs} \\
=& \sum_{t=1}^{T} \|\tilde{U}_{t-1}(I - W) + (-\hat{V}_{t-2} + \hat{V}_{t-1})W\|_{abs} \\
=& \sum_{t=1}^{T} \|\tilde{U}_{t-1}(QQ^T - Q\Lambda Q^T) + (-\hat{V}_{t-2} + \hat{V}_{t-1})W\|_{abs} \\
=& \sum_{t=1}^{T} \|\tilde{U}_{t-1}(\sum_{l=2}^{N} q_l(1-\lambda_l)q_l^T) + (-\hat{V}_{t-2} + \hat{V}_{t-1})W\|_{abs} \\
\leq& \sum_{t=1}^{T} \|\sum_{k=1}^{t-1}(-\hat{V}_{t-2-k} + \hat{V}_{t-1-k})\sum_{l=2}^{N} q_l\lambda_l^k(1-\lambda_l)q_l^T\|_{abs} + \sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})W\|_{abs} \\
\leq& \sum_{t=1}^{T}\left(\sum_{k=1}^{t-1}\| - \hat{V}_{t-2-k} + \hat{V}_{t-1-k}\|_{abs}\sqrt{N}\lambda^k\right) + \sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs} \\
=& \sum_{t=1}^{T}\left(\sum_{o=1}^{t-1}\| - \hat{V}_{o-2} + \hat{V}_{o-1}\|_{abs}\sqrt{N}\lambda^{t-o}\right) + \sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs}
\end{aligned}$$

$$
\begin{aligned}
&= \sum_{o=1}^{T-1} \sum_{t=o+1}^{T} \left(\| - \hat{V}_{o-2} + \hat{V}_{o-1} \|_{abs} \sqrt{N} \lambda^{t-o} \right) + \sum_{t=1}^{T} \| (-\hat{V}_{t-2} + \hat{V}_{t-1}) \|_{abs} \\
&\leq \sum_{o=1}^{T-1} \frac{\lambda}{1-\lambda} \left(\| - \hat{V}_{o-2} + \hat{V}_{o-1} \|_{abs} \sqrt{N} \right) + \sum_{t=1}^{T} \| (-\hat{V}_{t-2} + \hat{V}_{t-1}) \|_{abs} \\
&\leq \frac{1}{1-\lambda} \sum_{t=1}^{T} \| (-\hat{V}_{t-2} + \hat{V}_{t-1}) \|_{abs} \sqrt{N} \qquad (41)
\end{aligned}
$$

Combining (40) and (41), we have

$$
T_1 \leq G_\infty^2 \frac{1}{2\epsilon^{1.5}} \frac{1}{N} \mathbb{E} \left[\frac{1}{1-\lambda} \sum_{t=1}^{T} \| (-\hat{V}_{t-2} + \hat{V}_{t-1}) \|_{abs} \sqrt{N} \right] \qquad (42)
$$

What remains is to bound $\sum_{t=1}^{T} \mathbb{E} \left[\| Z_{t+1} - Z_t \|^2 \right]$. By update rule of Z_t, we have

$$
\begin{aligned}
&\| Z_{t+1} - Z_t \|^2 \\
&= \left\| \alpha \frac{\beta_1}{1-\beta_1} \frac{1}{N} \sum_{i=1}^{N} m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}} \right) - \alpha \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \\
&\leq 2\alpha^2 \left\| \frac{\beta_1}{1-\beta_1} \frac{1}{N} \sum_{i=1}^{N} m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}} \right) \right\|^2 + 2\alpha^2 \left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \\
&\leq 2\alpha^2 \left(\frac{\beta_1}{1-\beta_1} \right)^2 G_\infty^2 \frac{1}{N} \sum_{i=1}^{N} \sum_{j=1}^{d} \frac{1}{\sqrt{\epsilon}} \left| \frac{1}{\sqrt{[u_{t-1,i}]_j}} - \frac{1}{\sqrt{[u_{t,i}]_j}} \right| + 2\alpha^2 \left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \\
&\leq 2\alpha^2 \left(\frac{\beta_1}{1-\beta_1} \right)^2 G_\infty^2 \frac{1}{N} \sum_{i=1}^{N} \sum_{j=1}^{d} \frac{1}{\sqrt{\epsilon}} \left| \frac{[u_{t,i}]_j - [u_{t-1,i}]_j}{2\epsilon^{1.5}} \right| + 2\alpha^2 \left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \\
&\leq 2\alpha^2 \left(\frac{\beta_1}{1-\beta_1} \right)^2 G_\infty^2 \frac{1}{N} \sum_{i=1}^{N} \sum_{j=1}^{d} \frac{1}{2\epsilon^2} |[\tilde{u}_{t,i}]_j - [\tilde{u}_{t-1,i}]_j| + 2\alpha^2 \left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \\
&= 2\alpha^2 \left(\frac{\beta_1}{1-\beta_1} \right)^2 G_\infty^2 \frac{1}{N} \frac{1}{2\epsilon^2} \| \tilde{U}_t - \tilde{U}_{t-1} \|_{abs} + 2\alpha^2 \left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \qquad (43)
\end{aligned}
$$

where the last inequality is again due to the definition that $[\tilde{u}_{t,i}]_j = \max([u_{t,i}]_j, \epsilon)$ and the fact that $\max(\cdot, \epsilon)$ is 1-Lipschitz.

Then, we have

$$
\begin{aligned}
&\sum_{t=1}^{T} \mathbb{E}[\| Z_{t+1} - Z_t \|^2] \\
&\leq 2\alpha^2 \left(\frac{\beta_1}{1-\beta_1} \right)^2 G_\infty^2 \frac{1}{N} \frac{1}{2\epsilon^2} \mathbb{E} \left[\sum_{t=1}^{T} \| \tilde{U}_t - \tilde{U}_{t-1} \|_{abs} \right] + 2\alpha^2 \sum_{t=1}^{T} \mathbb{E} \left[\left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \right] \\
&\leq \alpha^2 \left(\frac{\beta_1}{1-\beta_1} \right)^2 \frac{G_\infty^2}{\sqrt{N}} \frac{1}{\epsilon^2} \frac{1}{1-\lambda} \mathbb{E} \left[\sum_{t=1}^{T} \| (-\hat{V}_{t-2} + \hat{V}_{t-1}) \|_{abs} \right] + 2\alpha^2 \sum_{t=1}^{T} \mathbb{E} \left[\left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \right] \qquad (44)
\end{aligned}
$$

where the last inequality is due to (41).

Now let's bound the last term on RHS of the above inequality. A trivial bound can be

$$
\sum_{t=1}^{T} \left\| \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \right\|^2 \leq \sum_{t=1}^{T} d G_\infty^2 \frac{1}{\epsilon}
$$

due to $\|g_{t,i}\| \leq G_\infty$ and $[u_{t,i}]_j \geq \epsilon, \forall j$ (this is easy to verify from update rule of $u_{t,i}$ and the assumption that $[v_{t,i}]_j \geq \epsilon, \forall i$). However, the above bound is independent of N, to get a better bound, we need a more involved analysis to show its dependency on N. To do this, we first notice that

$$\begin{aligned}
&\mathbb{E}_{G_t|G_{1:t-1}}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{g_{t,i}}{\sqrt{u_{t,i}}}\right\|^2\right] \\
=&\mathbb{E}_{G_t|G_{1:t-1}}\left[\frac{1}{N^2}\sum_{i=1}^N\sum_{j=1}^N\left\langle\frac{\nabla f_i(x_{t,i})+\xi_{t,i}}{\sqrt{u_{t,i}}}, \frac{\nabla f_j(x_{t,j})+\xi_{t,j}}{\sqrt{u_{t,j}}}\right\rangle\right] \\
\overset{(a)}{=}&\mathbb{E}_{G_t|G_{1:t-1}}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2\right] + \mathbb{E}_{G_t|G_{1:t-1}}\left[\frac{1}{N^2}\sum_{i=1}^N\left\|\frac{\xi_{t,i}}{\sqrt{u_{t,i}}}\right\|^2\right] \\
\overset{(b)}{=}&\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2 + \frac{1}{N^2}\sum_{i=1}^N\sum_{l=1}^d\frac{\mathbb{E}_{G_t|G_{1:t-1}}[[\xi_{t,i}]_l^2]}{[u_{t,i}]_l} \\
\overset{(c)}{\leq}&\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2 + \frac{d}{N}\frac{\sigma^2}{\epsilon}
\end{aligned} \tag{45}$$

where (a) is due to $\mathbb{E}_{G_t|G_{1:t-1}}[\xi_{t,i}] = 0$ and $\xi_{t,i}$ is independent of $x_{t,j}, \forall j$, $u_{t,j}, \forall j$, and $\xi_j, \forall j \neq i$, (b) comes from the fact that $x_{t,i}$, $u_{t,i}$ are fixed given $G_{1:t}$, (c) is due to $\mathbb{E}_{G_t|G_{1:t-1}}[[\xi_{t,i}]_l^2 \leq \sigma^2$ and $[u_{t.i}]_l \geq \epsilon$ by definition.

Then we have

$$\begin{aligned}
\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{g_{t,i}}{\sqrt{u_{t,i}}}\right\|^2\right] =&\mathbb{E}_{G_{1:t-1}}\left[\mathbb{E}_{G_t|G_{1:t-1}}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{g_{t,i}}{\sqrt{u_{t,i}}}\right\|^2\right]\right] \\
\leq&\mathbb{E}_{G_{1:t-1}}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2 + \frac{d}{N}\frac{\sigma^2}{\epsilon}\right] \\
=&\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2\right] + \frac{d}{N}\frac{\sigma^2}{\epsilon}
\end{aligned} \tag{46}$$

In traditional analysis of SGD-like distributed algorithms, the term corresponding to $\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2\right]$ will be merged with the first order descent when the stepsize is chosen to be small enough. However, in our case, the term cannot be merged because it is different from the first order descent in our algorithm. A brute-force upper bound is possible but this will lead to a worse convergence rate in terms of N. Thus, we need a more detailed analysis for the term in the following.

$$\begin{aligned}
\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2\right] =&\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}} + \frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i})\odot\left(\frac{1}{\sqrt{u_{t,i}}}-\frac{1}{\sqrt{\overline{U}_t}}\right)\right\|^2\right] \\
\leq&2\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right] + 2\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \nabla f_i(x_{t,i})\odot\left(\frac{1}{\sqrt{u_{t,i}}}-\frac{1}{\sqrt{\overline{U}_t}}\right)\right\|^2\right] \\
\leq&2\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^N \frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right] + 2\mathbb{E}\left[\frac{1}{N}\sum_{i=1}^N\left\|\nabla f_i(x_{t,i})\odot\left(\frac{1}{\sqrt{u_{t,i}}}-\frac{1}{\sqrt{\overline{U}_t}}\right)\right\|^2\right]
\end{aligned}$$

$$\leq 2\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right] + 2\mathbb{E}\left[\frac{1}{N}\sum_{i=1}^{N}G_\infty^2\frac{1}{\sqrt{\epsilon}}\left\|\frac{1}{\sqrt{u_{t,i}}}-\frac{1}{\sqrt{\overline{U}_t}}\right\|_1\right] \tag{47}$$

Summing over T, we have

$$\begin{aligned}
&\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{u_{t,i}}}\right\|^2\right] \\
\leq& 2\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right] + 2\sum_{t=1}^{T}\mathbb{E}\left[\frac{1}{N}\sum_{i=1}^{N}G_\infty^2\frac{1}{\sqrt{\epsilon}}\left\|\frac{1}{\sqrt{u_{t,i}}}-\frac{1}{\sqrt{\overline{U}_t}}\right\|_1\right]
\end{aligned} \tag{48}$$

For the last term on RHS of (48), we can bound it similarly as what we did for T_2 from (32) to (38), which yields

$$\begin{aligned}
&\sum_{t=1}^{T}\mathbb{E}\left[\frac{1}{N}\sum_{i=1}^{N}G_\infty^2\frac{1}{\sqrt{\epsilon}}\left\|\frac{1}{\sqrt{u_{t,i}}}-\frac{1}{\sqrt{\overline{U}_t}}\right\|_1\right] \\
\leq& \sum_{t=1}^{T}\mathbb{E}\left[\frac{1}{N}\sum_{i=1}^{N}G_\infty^2\frac{1}{\sqrt{\epsilon}}\frac{1}{2\epsilon^{1.5}}\left\|u_{t,i}-\overline{U}_t\right\|_1\right] \\
=& \sum_{t=1}^{T}\mathbb{E}\left[\frac{1}{N}G_\infty^2\frac{1}{2\epsilon^2}\left\|\overline{U}_t\mathbf{1}^T-U_t\right\|_{abs}\right] \\
\leq& \sum_{t=1}^{T}\mathbb{E}\left[\frac{1}{N}G_\infty^2\frac{1}{2\epsilon^2}\|-\sum_{l=2}^{N}\tilde{U}_t q_l q_l^T\|_{abs}\right] \\
\leq& \frac{1}{\sqrt{N}}G_\infty^2\frac{1}{2\epsilon^2}\mathbb{E}\left[\sum_{o=0}^{T-1}\frac{\lambda}{1-\lambda}\|(-\hat{V}_{o-1}+\hat{V}_o)\|_{abs}\right]
\end{aligned} \tag{49}$$

Further, we have

$$\begin{aligned}
&\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right] \\
\leq& 2\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(\overline{X}_t)}{\sqrt{\overline{U}_t}}\right\|^2\right] + 2\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(\overline{X}_t)-\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right] \\
=& 2\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\sqrt{\overline{U}_t}}\right\|^2\right] + 2\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(\overline{X}_t)-\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right]
\end{aligned} \tag{50}$$

and the last term on RHS of the above inequality can be bounded following similar procedures from (25) to (30), as what we did for T_3. Completing the procedures yields

$$\begin{aligned}
&\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{\nabla f_i(\overline{X}_t)-\nabla f_i(x_{t,i})}{\sqrt{\overline{U}_t}}\right\|^2\right] \\
\leq& \sum_{t=1}^{T}\mathbb{E}\left[L\frac{1}{\epsilon}\frac{1}{N}\sum_{i=1}^{N}\left\|x_{t,i}-\overline{X}_t\right\|^2\right] \\
\leq& \sum_{t=1}^{T}\mathbb{E}\left[L\frac{1}{\epsilon}\frac{1}{N}\alpha^2\left(\frac{1}{1-\lambda}\right)NdG_\infty^2\frac{1}{\epsilon}\right] \\
=& TL\frac{1}{\epsilon^2}\alpha^2\left(\frac{1}{1-\lambda}\right)dG_\infty^2
\end{aligned} \tag{51}$$

Finally, combining (46) to (51), we get

$$
\begin{aligned}
&\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{1}{N}\sum_{i=1}^{N}\frac{g_{t,i}}{\sqrt{u_{t,i}}}\right\|^2\right]\\
\leq& 4\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\sqrt{\overline{U}_t}}\right\|^2\right]+4TL\frac{1}{\epsilon^2}\alpha^2\left(\frac{1}{1-\lambda}\right)dG_\infty^2\\
&+2\frac{1}{\sqrt{N}}G_\infty^2\frac{1}{2\epsilon^2}\mathbb{E}\left[\sum_{o=0}^{T-1}\frac{\lambda}{1-\lambda}\|(-\hat{V}_{o-1}+\hat{V}_o)\|_{abs}\right]+T\frac{d}{N}\frac{\sigma^2}{\epsilon}\\
\leq& 4\frac{1}{\sqrt{\epsilon}}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right]+4TL\frac{1}{\epsilon^2}\alpha^2\left(\frac{1}{1-\lambda}\right)dG_\infty^2\\
&+2\frac{1}{\sqrt{N}}G_\infty^2\frac{1}{2\epsilon^2}\mathbb{E}\left[\sum_{o=0}^{T-1}\frac{\lambda}{1-\lambda}\|(-\hat{V}_{o-1}+\hat{V}_o)\|_{abs}\right]+T\frac{d}{N}\frac{\sigma^2}{\epsilon}. \qquad (52)
\end{aligned}
$$

where the last inequality is due to each element of $\overline{U}_t$ is lower bounded by ϵ by definition.

Combining all above, we can have

$$
\begin{aligned}
&\frac{1}{T}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right]\\
\leq&\frac{2}{T\alpha}(\mathbb{E}[f(Z_1)]-\mathbb{E}[f(Z_{T+1})])\\
&+\frac{L}{T}\alpha\left(\frac{\beta_1}{1-\beta_1}\right)^2\frac{G_\infty^2}{\sqrt{N}}\frac{1}{\epsilon^2}\frac{1}{1-\lambda}\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2}+\hat{V}_{t-1})\|_{abs}\right]\\
&+\frac{8L}{T}\alpha\frac{1}{\sqrt{\epsilon}}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right]+8L^2\alpha\frac{1}{\epsilon^2}\alpha^2\left(\frac{1}{1-\lambda}\right)dG_\infty^2\\
&+\frac{4L}{T}\alpha\frac{1}{\sqrt{N}}G_\infty^2\frac{1}{2\epsilon^2}\mathbb{E}\left[\sum_{o=0}^{T-1}\frac{\lambda}{1-\lambda}\|(-\hat{V}_{o-1}+\hat{V}_o)\|_{abs}\right]+2L\alpha\frac{d}{N}\frac{\sigma^2}{\epsilon}\\
&+\frac{2}{T}\frac{\beta_1}{1-\beta_1}G_\infty^2\frac{1}{2\epsilon^{1.5}}\frac{1}{\sqrt{N}}\mathbb{E}\left[\frac{1}{1-\lambda}\sum_{t=1}^{T}\|(-\hat{V}_{t-2}+\hat{V}_{t-1})\|_{abs}\right]\\
&+\frac{2}{T}\frac{G_\infty^2}{\sqrt{N}}\frac{1}{2\epsilon^{1.5}}\frac{\lambda}{1-\lambda}\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2}+\hat{V}_{t-1})\|_{abs}\right]\\
&+\frac{3}{T}\left(\sum_{t=1}^{T}L\left(\frac{1}{1-\lambda}\right)^2\alpha^2dG_\infty^2\frac{1}{\epsilon^{1.5}}+\sum_{t=1}^{T}L\left(\frac{\beta_1}{1-\beta_1}\right)^2\alpha^2d\frac{G_\infty^2}{\epsilon^{1.5}}\right)\\
=&\frac{2}{T\alpha}(\mathbb{E}[f(Z_1)]-\mathbb{E}[f(Z_{T+1})])+2L\alpha\frac{d}{N}\frac{\sigma^2}{\epsilon}+8L\alpha\frac{1}{\sqrt{\epsilon}}\frac{1}{T}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right]\\
&+3\alpha^2d\left(\left(\frac{\beta_1}{1-\beta_1}\right)^2+\left(\frac{1}{1-\lambda}\right)^2\right)L\frac{G_\infty^2}{\epsilon^{1.5}}+8\alpha^3L^2\left(\frac{1}{1-\lambda}\right)d\frac{G_\infty^2}{\epsilon^2}\\
&+\frac{1}{T\epsilon^{1.5}}\frac{G_\infty^2}{\sqrt{N}}\frac{1}{1-\lambda}\left(L\alpha\left(\frac{\beta_1}{1-\beta_1}\right)^2\frac{1}{\epsilon^{0.5}}+\lambda+\frac{\beta_1}{1-\beta_1}+2L\alpha\frac{1}{\epsilon^{0.5}}\lambda\right)\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2}+\hat{V}_{t-1})\|_{abs}\right]. \qquad (53)
\end{aligned}
$$

Set $\alpha = \frac{1}{\sqrt{dT}}$ and when $\alpha \leq \frac{\epsilon^{0.5}}{16L}$, we further have

$$
\begin{aligned}
&\frac{1}{T}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right] \\
\leq&\frac{4}{T\alpha}(\mathbb{E}[f(Z_1)] - \mathbb{E}[f(Z_{T+1})]) + 4L\alpha\frac{d}{N}\frac{\sigma^2}{\epsilon} \\
&+ 6\alpha^2 d\left(\left(\frac{\beta_1}{1-\beta_1}\right)^2 + \left(\frac{1}{1-\lambda}\right)^2\right) L\frac{G_\infty^2}{\epsilon^{1.5}} + 16\alpha^3 L^2\left(\frac{1}{1-\lambda}\right) d\frac{G_\infty^2}{\epsilon^2} \\
&+ \frac{2}{T\epsilon^{1.5}}\frac{G_\infty^2}{\sqrt{N}}\frac{1}{1-\lambda}\left(L\alpha\left(\frac{\beta_1}{1-\beta_1}\right)^2\frac{1}{\epsilon^{0.5}} + \lambda + \frac{\beta_1}{1-\beta_1} + 2L\alpha\frac{1}{\epsilon^{0.5}}\lambda\right)\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs}\right] \\
=&\frac{4\sqrt{d}}{\sqrt{T}}(\mathbb{E}[f(Z_1)] - \mathbb{E}[f(Z_{T+1})]) + 4L\frac{\sqrt{d}}{\sqrt{T}}\frac{1}{N}\frac{\sigma^2}{\epsilon} \\
&+ 6\frac{1}{T}\left(\left(\frac{\beta_1}{1-\beta_1}\right)^2 + \left(\frac{1}{1-\lambda}\right)^2\right) L\frac{G_\infty^2}{\epsilon^{1.5}} + 16\frac{1}{T^{1.5}d^{0.5}}L^2\left(\frac{1}{1-\lambda}\right)\frac{G_\infty^2}{\epsilon^2} \\
&+ \frac{2}{T\epsilon^{1.5}}\frac{G_\infty^2}{\sqrt{N}}\frac{1}{1-\lambda}\left(\frac{L}{\sqrt{Td}}\left(\frac{\beta_1}{1-\beta_1}\right)^2\frac{1}{\epsilon^{0.5}} + \lambda + \frac{\beta_1}{1-\beta_1} + 2\frac{L}{\sqrt{Td}}\frac{1}{\epsilon^{0.5}}\lambda\right)\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs}\right] \\
\leq&C_1\frac{\sqrt{d}}{\sqrt{T}}\left(\mathbb{E}[f(Z_1)] - \min_z f(z) + \frac{\sigma^2}{N}\right) + \frac{1}{T}C_2 + \frac{1}{T^{1.5}d^{0.5}}C_3 \\
&+ \left(\frac{1}{TN^{0.5}}C_4 + \frac{1}{T^{1.5}d^{0.5}N^{0.5}}C_5\right)\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs}\right] \qquad (54)
\end{aligned}
$$

where the first inequality is obtained by moving the term $8L\alpha\frac{1}{\sqrt{\epsilon}}\frac{1}{T}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right]$ on the RHS of (53) to the LHS to cancel it using the assumption $8L\alpha\frac{1}{\sqrt{\epsilon}} \leq \frac{1}{2}$ followed by multiplying both sides by 2, and the constants introduced in the last step are defined as following

$$
\begin{aligned}
C_1 =& \max(4, 4L/\epsilon) \\
C_2 =& 6\left(\left(\frac{\beta_1}{1-\beta_1}\right)^2 + \left(\frac{1}{1-\lambda}\right)^2\right) L\frac{G_\infty^2}{\epsilon^{1.5}} \\
C_3 =& 16L^2\left(\frac{1}{1-\lambda}\right)\frac{G_\infty^2}{\epsilon^2} \\
C_4 =& \frac{2}{\epsilon^{1.5}}\frac{1}{1-\lambda}\left(\lambda + \frac{\beta_1}{1-\beta_1}\right)G_\infty^2 \\
C_5 =& \frac{2}{\epsilon^2}\frac{1}{1-\lambda}L\left(\frac{\beta_1}{1-\beta_1}\right)^2 G_\infty^2 + \frac{4}{\epsilon^2}\frac{\lambda}{1-\lambda}LG_\infty^2. \qquad (55)
\end{aligned}
$$

Substituting into $Z_1 = \overline{X}_1$ completes the proof □

A.2 Proof of Theorem 3

By Theorem 2, we know under the assumptions of the theorem, we have

$$
\begin{aligned}
\frac{1}{T}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right] \leq& C_1\frac{\sqrt{d}}{\sqrt{T}}\left(\mathbb{E}[f(\overline{X}_1)] - \min_z f(z)] + \frac{\sigma^2}{N}\right) + \frac{1}{T}C_2 + \frac{1}{T^{1.5}d^{0.5}}C_3 \\
&+ \left(\frac{1}{TN^{0.5}}C_4 + \frac{1}{T^{1.5}d^{0.5}N^{0.5}}C_5\right)\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2} + \hat{V}_{t-1})\|_{abs}\right] \qquad (56)
\end{aligned}
$$

where $\|\cdot\|_{abs}$ denotes the entry-wise L_1 norm of a matrix (i.e $\|A\|_{abs} = \sum_{i,j}|A_{ij}|$) and C_1, C_2, C_3, C_4, C_5 are defined in Theorem 2.

Since Algorithm 3 is a special case of 2, building on result of Theorem 2, we just need to characterize the growth speed of $\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2}+\hat{V}_{t-1})\|_{abs}\right]$ to prove convergence of Algorithm 3. By the update rule of Algorithm 3, we know $\hat{V}_t$ is non decreasing and thus

$$\begin{aligned}
&\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2}+\hat{V}_{t-1})\|_{abs}\right]\\
=&\mathbb{E}\left[\sum_{t=1}^{T}\sum_{i=1}^{N}\sum_{j=1}^{d}|-[\hat{v}_{t-2,i}]_j+[\hat{v}_{t-1,i}]_j|\right]\\
=&\mathbb{E}\left[\sum_{t=1}^{T}\sum_{i=1}^{N}\sum_{j=1}^{d}(-[\hat{v}_{t-2,i}]_j+[\hat{v}_{t-1,i}]_j)\right]\\
=&\mathbb{E}\left[\sum_{i=1}^{N}\sum_{j=1}^{d}(-[\hat{v}_{-1,i}]_j+[\hat{v}_{T-1,i}]_j)\right]\\
=&\mathbb{E}\left[\sum_{i=1}^{N}\sum_{j=1}^{d}(-[\hat{v}_{0,i}]_j+[\hat{v}_{T-1,i}]_j)\right]
\end{aligned} \tag{57}$$

where the last equality is because we defined $\hat{V}_{-1} \triangleq \hat{V}_0$ previously.

Further, because $\|g_{t,i}\|_\infty \le G_\infty, \forall t, i$ and $v_{t,i}$ is a exponential moving average of $g_{k,i}^2, k = 1, 2, \cdots, t$, we know $|[v_{t,i}]_j| \le G_\infty^2, \forall t, i, j$. In addition, by update rule of $\hat{V}_t$, we also know each element of $\hat{V}_t$ also cannot be greater than G_∞^2, i.e. $|[\hat{v}_{t,i}]_j| \le G_\infty^2, \forall t, i, j$.

Given the fact that $[\hat{v}_{0,i}]_j \ge 0$, we have

$$\mathbb{E}\left[\sum_{t=1}^{T}\|(-\hat{V}_{t-2}+\hat{V}_{t-1})\|_{abs}\right] = \mathbb{E}\left[\sum_{i=1}^{N}\sum_{j=1}^{d}(-[\hat{v}_{0,i}]_j+[\hat{v}_{T-1,i}]_j)\right] \le \mathbb{E}\left[\sum_{i=1}^{N}\sum_{j=1}^{d}G_\infty^2\right] = NdG_\infty^2$$

Substituting the above into (56), we have

$$\begin{aligned}
\frac{1}{T}\sum_{t=1}^{T}\mathbb{E}\left[\left\|\frac{\nabla f(\overline{X}_t)}{\overline{U}_t^{1/4}}\right\|^2\right] \le& C_1\frac{\sqrt{d}}{\sqrt{T}}\left(\mathbb{E}[f(\overline{X}_1)]-\min_z f(z)+\frac{\sigma^2}{N}\right)+\frac{1}{T}C_2+\frac{1}{T^{1.5}d^{0.5}}C_3\\
&+\frac{d}{T}C_4\sqrt{N}G_\infty^2+\frac{\sqrt{d}}{T^{1.5}}C_5\sqrt{N}G_\infty^2\\
=&C_1'\frac{\sqrt{d}}{\sqrt{T}}\left(\mathbb{E}[f(\overline{X}_1)]-\min_z f(z)+\frac{\sigma^2}{N}\right)+\frac{1}{T}C_2'+\frac{1}{T^{1.5}d^{0.5}}C_3'\\
&+\frac{d}{T}\sqrt{N}C_4'+\frac{\sqrt{d}}{T^{1.5}}\sqrt{N}C_5'
\end{aligned} \tag{58}$$

where we have

$$\begin{aligned}
C_1' &= C_1\\
C_2' &= C_2\\
C_3' &= C_3\\
C_4' &= C_4G_\infty^2\\
C_5' &= C_5G_\infty^2
\end{aligned} \tag{59}$$

The proof is complete. □

A.3 Proof of Lemmas

Lemma 1. *For the sequence defined in (14), we have*

$$Z_{t+1} - Z_t = \alpha \frac{\beta_1}{1-\beta_1} \frac{1}{N} \sum_{i=1}^{N} m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}}\right) - \alpha \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}} \tag{60}$$

Proof: By update rule of Algorithm 2, we first have

$$\begin{aligned}
\overline{X}_{t+1} =& \frac{1}{N} \sum_{i=1}^{N} x_{t+1,i} \\
=& \frac{1}{N} \sum_{i=1}^{N} \left(x_{t+0.5,i} - \alpha \frac{m_{t,i}}{\sqrt{u_{t,i}}} \right) \\
=& \frac{1}{N} \sum_{i=1}^{N} \left(\sum_{j=1}^{N} W_{ij} x_{t,j} - \alpha \frac{m_{t,i}}{\sqrt{u_{t,i}}} \right) \\
\overset{(i)}{=}& \left(\frac{1}{N} \sum_{j=1}^{N} x_{t,j} \right) - \frac{1}{N} \sum_{i=1}^{N} \alpha \frac{m_{t,i}}{\sqrt{u_{t,i}}} \\
=& \overline{X}_t - \frac{1}{N} \sum_{i=1}^{N} \alpha \frac{m_{t,i}}{\sqrt{u_{t,i}}}
\end{aligned} \tag{61}$$

where (i) is due to an interchange of summation and $\sum_{i=1} W_{ij} = 1$.

Then, we have

$$\begin{aligned}
Z_{t+1} - Z_t =& \overline{X}_{t+1} - \overline{X}_t + \frac{\beta_1}{1-\beta_1}(\overline{X}_{t+1} - \overline{X}_t) - \frac{\beta_1}{1-\beta_1}(\overline{X}_{t+1} - \overline{X}_t) \\
=& \frac{1}{1-\beta_1}(\overline{X}_{t+1} - \overline{X}_t) - \frac{\beta_1}{1-\beta_1}(\overline{X}_{t+1} - \overline{X}_t) \\
=& \frac{1}{1-\beta_1} \left(-\frac{1}{N} \sum_{i=1}^{N} \alpha \frac{m_{t,i}}{\sqrt{u_{t,i}}} \right) - \frac{\beta_1}{1-\beta_1} \left(-\frac{1}{N} \sum_{i=1}^{N} \alpha \frac{m_{t-1,i}}{\sqrt{u_{t-1,i}}} \right) \\
=& \frac{1}{1-\beta_1} \left(-\frac{1}{N} \sum_{i=1}^{N} \alpha \frac{\beta_1 m_{t-1,i} + (1-\beta_1) g_{t,i}}{\sqrt{u_{t,i}}} \right) - \frac{\beta_1}{1-\beta_1} \left(-\frac{1}{N} \sum_{i=1}^{N} \alpha \frac{m_{t-1,i}}{\sqrt{u_{t-1,i}}} \right) \\
=& \alpha \frac{\beta_1}{1-\beta_1} \frac{1}{N} \sum_{i=1}^{N} m_{t-1,i} \odot \left(\frac{1}{\sqrt{u_{t-1,i}}} - \frac{1}{\sqrt{u_{t,i}}} \right) - \alpha \frac{1}{N} \sum_{i=1}^{N} \frac{g_{t,i}}{\sqrt{u_{t,i}}}
\end{aligned} \tag{62}$$

which is the desired result. □

Lemma 2. *Given a set of numbers $a_1, \cdots, a_n$ and denote their mean to be $\bar{a} = \frac{1}{n}\sum_{i=1}^{n} a_i$. In addition, define $b_i(r) \triangleq= \max(a_i, r)$ and $\bar{b}(r) = \frac{1}{n}\sum_{i=1}^{n} b_i(r)$. For any r and r' with $r' \geq r$ we have*

$$\sum_{i=1}^{n} |b_i(r) - \bar{b}(r)| \geq \sum_{i=1}^{n} |b_i(r') - \bar{b}(r')| \tag{63}$$

and when $r \leq \min_{i\in[n]} a_i$, we have

$$\sum_{i=1}^{n} |b_i(r) - \bar{b}(r)| = \sum_{i=1}^{n} |a_i - \bar{a}| \tag{64}$$

Proof: Without loss of generality, let's assume $a_i \leq a_j$ when $i < j$, i.e. a_i is a non-decreasing sequence. Define

$$h(r) = \sum_{i=1}^{n} |b_i(r) - \bar{b}(r)| = \sum_{i=1}^{n} |\max(a_i, r) - \frac{1}{n}\sum_{j=1}^{n} \max(a_j, r)|, \tag{65}$$

we need to prove that h is a non-increasing function of r. First, it is easy to see that h is a continuous function of r with non-differentiable points $r = a_i, i \in [n]$, thus h is a piece-wise linear function.

Next, we will prove that $h(r)$ is non-increasing in each piece. Define $l(r)$ to be the largest index with $a(l(r)) < r$, and $s(r)$ to be the largest index with $a_{s(r)} < \bar{b}(r)$. Note that we have $b_i(r) = r, \forall i \leq l(r)$ and $b_i(r) - \bar{b}(r) \leq 0, \forall i \leq s(r)$ because a_i is a non-decreasing sequence. Therefore, we have

$$h(r) = \sum_{i=1}^{l(r)} (\bar{b}(r) - r) + \sum_{i=l(r)+1}^{s(r)} (\bar{b}(r) - a_i) + \sum_{i=s(r)+1}^{n} (a_i - \bar{b}(r)). \tag{66}$$

and

$$\bar{b}(r) = \frac{1}{n} \left(l(r)r + \sum_{i=l(r)+1}^{n} a_i \right) \tag{67}$$

Taking derivative of the above form, we know the derivative of $h(r)$ at differentiable points is

$$\begin{aligned} h'(r) =& l(r)(\frac{l(r)}{n} - 1) + (s(r) - l(r))\frac{l(r)}{n} - (n - s(r))\frac{l(r)}{n} \\ =& \frac{l(r)}{n}((l(r) - n) + (s(r) - l(r)) - (n - s(r))) \end{aligned} \tag{68}$$

Since we have $s(r) \leq n$ we know $(l(r) - n) + (s(r) - l(r)) - (n - s(r)) \leq 0$ and thus

$$h'(r) \leq 0 \tag{69}$$

which means $h(r)$ is non-increasing in each piece. Combining with the fact that $h(r)$ is continuous, (64) is proven.

When $r \leq a(i)$, we have $b(i) = \max(a_i, r) = r, \forall r \in [n]$ and $\bar{b}(r) = \frac{1}{n}\sum_{i=1}^{n} a_i = \bar{a}$ which proves (65). □

A.4 Additional experiments and details

In this section, we compare the learning curves of different algorithms with different stepsizes on heterogeneous data distribution. We use 5 nodes and the heterogeneous data distribution is created by assigning each node with data of only two labels and there are no overlapping labels between different nodes. For all algorithms, we compare stepsizes in the set [1e-1, 1e-2, 1e-3, 1e-4, 1e-5, 1e-6].

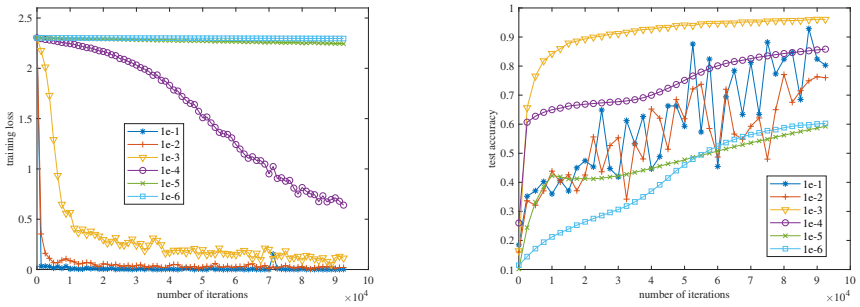

(a) Training loss (b) Test accuracy

Figure 2: Performance comparison of different stepsizes for DGD

Figure 2 shows the training loss and test accuracy of DGD, it can be seen that the stepsize 1e-3 works best for DGD in terms of test accuracy and 1e-1 works best in terms of training loss. The difference is caused by the inconsistency among the value of parameters on different nodes when the stepsize is large. The training loss is calculated as the average of the loss value of different local models evaluated on their local training batch. Thus, though the training loss is small evaluated at a particular

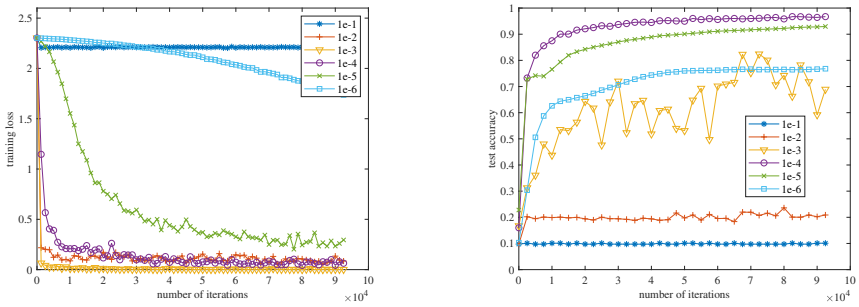

(a) Training loss (b) Test accuracy

Figure 3: Performance comparison of different stepsizes for decentralized AMSGrad

node, the test accuracy will be low when evaluating data with labels not seen by the node (recall that each node contains data with different labels).

Figure 3 shows the performance of decentralized AMSGrad with different stepsizes, we can see its best performance is better than DGD and the performance is stabler (the test performance is less sensitive to stepsize choice).

Figure 4 shows the performance of DADM, as it can be expected, the performance of DADAM is not as good as DGD and decentralized AMSGrad since it is not a convergent algorithm and the heterogeneity in data amplified the non-convergence issue of DADAM.

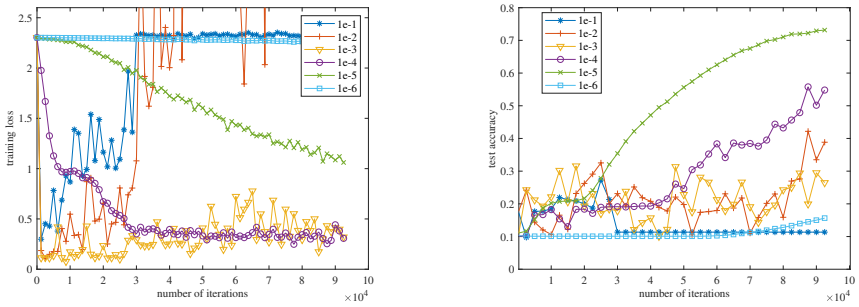

(a) Training loss (b) Test accuracy

Figure 4: Performance comparison of different stepsizes for DADAM

From the experiments above, we can see the advantages of decentralized AMSGrad in terms of both performance and ease of parameter tuning, and the importance of ensuring the theoretical convergence of algorithms.